

Data Mining with Algorithms – Random Forest and Support Vector Machine

Alex C. Berger

New Jersey Institute of Technology

Ying Wu College of Computing

November 16th, 2024

Table of Contents

Title Page	1
Table of Contents	2
Introduction	3
Datasets	4-8
Algorithms	9-16
Methodology	17-18
Implementation	19-23
Results	24-29
Discussion	30-31
Conclusion	32
Appendix	33-85

1. Introduction

Overview

In this project, we aim to implement and evaluate the performance of two distinct classification algorithms selected from different categories. These algorithms will be applied to three data sets to compare their accuracy, robustness, and adaptability across varying data distributions and characteristics. The results will be analyzed using standardized evaluation metrics, including accuracy, 10-fold cross-validation, and ROC/AUC where applicable.

Project Background

This project builds upon the principles and methodologies learned in the CS 634 Data Mining course. It provides an opportunity to apply theoretical concepts to practical scenarios by selecting appropriate algorithms, preprocessing data, and interpreting the results of experiments. The focus is on understanding how the chosen algorithms perform on diverse datasets and how the two models compare on these datasets.

Project Focus

1. **Algorithm Performance Analysis:** Compare the classification accuracy of two algorithms from different categories when applied to a primary dataset.
2. **Generalizability Assessment:** Extend the evaluation to two additional datasets to assess the robustness and adaptability of the algorithms.

Through these analyses, this project seeks to identify key strengths and limitations of the selected algorithms. This exploration will also highlight the impact of dataset characteristics on model performance.

2. Datasets

2.1 Dataset 1: Mobile Device Usage and User Behavior Dataset

- **Source:** Kaggle
- **Description:** This dataset contains behavioral data on mobile device usage. It includes information such as screen-on time, battery drain, app usage, age, and gender of users.
- **Size:**
 - Number of Records: 700
 - Number of Features: 11
- **Target Variable:** User Behavior Class (Multi-class: 1, 2, 3, 4, 5).
- **Preprocessing Steps:**
 - Applied one-hot encoding to categorical features.
 - Standardized numerical features for algorithm compatibility.

2.2 Dataset 2: Loan Approval Dataset

- **Source:** Kaggle
- **Description:** This dataset includes financial and demographic information used to determine loan approval. Features include applicant income, education level, and self-employment status.
- **Size:**
 - Number of Records: 4,269
 - Number of Features: 13
- **Target Variable:** Loan Status (binary: Approved/Rejected).
- **Preprocessing Steps:**
 - Created a binary target column (loan_status_binary).
 - One-hot encoded categorical features (e.g., education, self-employed).
 - Standardized numerical features for consistency.

2.3 Dataset 3: Weather Classification Dataset

- **Source:** Kaggle
- **Description:** This dataset includes weather-related data such as temperature, humidity, cloud cover, and season to classify weather types (e.g., Sunny, Rainy).
- **Size:**
 - Number of Records: 2500 (sampled from the original dataset for time issues)

- Number of Features: 11
- **Target Variable:** Weather Type (Multi-class: Sunny, Cloudy, Rainy).
- **Preprocessing Steps:**
 - Applied one-hot encoding to categorical features (e.g., cloud cover, season).
 - Standardized numerical features for algorithmic compatibility.
 - Created a binary target column (target) for classification.

2.4 Screenshots

Dataset #1

	User ID	Device Model	Operating System	App Usage Time (min/day)	Screen On Time (hours/day)	Battery Drain (mAh/day)	Number of Apps Installed	Data Usage (MB/day)	Age	Gender	User Behavior Class
0	1	Google Pixel 5	Android	393	6.4	1872	67	1122	40	Male	4
1	2	OnePlus 9	Android	268	4.7	1331	42	944	47	Female	3
2	3	Xiaomi Mi 11	Android	154	4.0	761	32	322	42	Male	2
3	4	Google Pixel 5	Android	239	4.8	1676	56	871	20	Male	3
4	5	iPhone 12	iOS	187	4.3	1367	58	988	31	Female	3

```

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 700 entries, 0 to 699
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User ID                               700 non-null   int64
1   Device Model                           700 non-null   object
2   Operating System                       700 non-null   object
3   App Usage Time (min/day)               700 non-null   int64
4   Screen On Time (hours/day)             700 non-null   float64
5   Battery Drain (mAh/day)                700 non-null   int64
6   Number of Apps Installed                700 non-null   int64
7   Data Usage (MB/day)                   700 non-null   int64
8   Age                                    700 non-null   int64
9   Gender                                 700 non-null   object
10  User Behavior Class                    700 non-null   int64
dtypes: float64(1), int64(7), object(3)
memory usage: 60.3+ KB

```

```

Missing Values:
User ID                0
Device Model           0
Operating System       0
App Usage Time (min/day)  0
Screen On Time (hours/day)  0
Battery Drain (mAh/day)  0
Number of Apps Installed  0
Data Usage (MB/day)     0
Age                    0
Gender                 0
User Behavior Class     0
dtype: int64

```

Statistical Summary:

	User ID	Device Model	Operating System	App Usage Time (min/day)	Screen On Time (hours/day)	Battery Drain (mAh/day)	Number of Apps Installed	Data Usage (MB/day)	Age	Gender	User Behavior Class
count	700.00000	700	700	700.000000	700.000000	700.000000	700.000000	700.000000	700.000000	700	700.000000
unique	NaN	5	2	NaN	NaN	NaN	NaN	NaN	NaN	2	NaN
top	NaN	Xiaomi Mi 11	Android	NaN	NaN	NaN	NaN	NaN	NaN	Male	NaN
freq	NaN	146	554	NaN	NaN	NaN	NaN	NaN	NaN	364	NaN
mean	350.50000	NaN	NaN	271.128571	5.272714	1525.158571	50.681429	929.742857	38.482857	NaN	2.990000
std	202.21688	NaN	NaN	177.199484	3.068584	819.136414	26.943324	640.451729	12.012916	NaN	1.401476
min	1.00000	NaN	NaN	30.000000	1.000000	302.000000	10.000000	102.000000	18.000000	NaN	1.000000
25%	175.75000	NaN	NaN	113.250000	2.500000	722.250000	26.000000	373.000000	28.000000	NaN	2.000000
50%	350.50000	NaN	NaN	227.500000	4.900000	1502.500000	49.000000	823.500000	38.000000	NaN	3.000000
75%	525.25000	NaN	NaN	434.250000	7.400000	2229.500000	74.000000	1341.000000	49.000000	NaN	4.000000
max	700.00000	NaN	NaN	598.000000	12.000000	2993.000000	99.000000	2497.000000	59.000000	NaN	5.000000

Dataset #2

	loan_id	no_of_dependents	education	self_employed	income_annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value
0	1	2	Graduate	No	9600000	29900000	12	778	2400000	17600000
1	2	0	Not Graduate	Yes	4100000	12200000	8	417	2700000	2200000
2	3	3	Graduate	No	9100000	29700000	20	506	7100000	4500000
3	4	3	Graduate	No	8200000	30700000	8	467	18200000	3300000
4	5	5	Not Graduate	Yes	9800000	24200000	20	382	12400000	8200000

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4269 entries, 0 to 4268

Data columns (total 13 columns):

#	Column	Non-Null	Count	Dtype
0	loan_id	4269	non-null	int64
1	no_of_dependents	4269	non-null	int64
2	education	4269	non-null	object
3	self_employed	4269	non-null	object
4	income_annum	4269	non-null	int64
5	loan_amount	4269	non-null	int64
6	loan_term	4269	non-null	int64
7	cibil_score	4269	non-null	int64
8	residential_assets_value	4269	non-null	int64
9	commercial_assets_value	4269	non-null	int64
10	luxury_assets_value	4269	non-null	int64
11	bank_asset_value	4269	non-null	int64
12	loan_status	4269	non-null	object

dtypes: int64(10), object(3)

memory usage: 433.7+ KB

	income_annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value	luxury_assets_value	bank_asset_value	loan_status
No	9600000	29900000	12	778	2400000	17600000	22700000	8000000	Approved
Yes	4100000	12200000	8	417	2700000	2200000	8800000	3300000	Rejected
No	9100000	29700000	20	506	7100000	4500000	33300000	12800000	Rejected
No	8200000	30700000	8	467	18200000	3300000	23300000	7900000	Rejected
Yes	9800000	24200000	20	382	12400000	8200000	29400000	5000000	Rejected

```

Missing Values:
loan_id          0
no_of_dependents 0
education        0
self_employed    0
income_annum     0
loan_amount      0
loan_term        0
cibil_score      0
residential_assets_value 0
commercial_assets_value 0
luxury_assets_value 0
bank_asset_value 0
loan_status      0
dtype: int64

```

Statistical Summary:

	loan_id	no_of_dependents	education	self_employed	income_annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value
count	4269.000000	4269.000000	4269	4269	4.269000e+03	4.269000e+03	4269.000000	4269.000000	4.269000e+03	
unique	NaN	NaN	2	2	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	Graduate	Yes	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	2144	2150	NaN	NaN	NaN	NaN	NaN	
mean	2135.000000	2.498712	NaN	NaN	5.059124e+06	1.513345e+07	10.900445	599.936051	7.472617e+06	
std	1232.498479	1.695910	NaN	NaN	2.806840e+06	9.043363e+06	5.709187	172.430401	6.503637e+06	
min	1.000000	0.000000	NaN	NaN	2.000000e+05	3.000000e+05	2.000000	300.000000	-1.000000e+05	
25%	1068.000000	1.000000	NaN	NaN	2.700000e+06	7.700000e+06	6.000000	453.000000	2.200000e+06	
50%	2135.000000	3.000000	NaN	NaN	5.100000e+06	1.450000e+07	10.000000	600.000000	5.600000e+06	
75%	3202.000000	4.000000	NaN	NaN	7.500000e+06	2.150000e+07	16.000000	748.000000	1.130000e+07	
max	4269.000000	5.000000	NaN	NaN	9.900000e+06	3.950000e+07	20.000000	900.000000	2.910000e+07	

income_annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value	luxury_assets_value	bank_asset_value	loan_status
4.269000e+03	4.269000e+03	4269.000000	4269.000000	4.269000e+03	4.269000e+03	4.269000e+03	4.269000e+03	4269
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	2
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Approved
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	2656
5.059124e+06	1.513345e+07	10.900445	599.936051	7.472617e+06	4.973155e+06	1.512631e+07	4.976692e+06	NaN
2.806840e+06	9.043363e+06	5.709187	172.430401	6.503637e+06	4.388966e+06	9.103754e+06	3.250185e+06	NaN
2.000000e+05	3.000000e+05	2.000000	300.000000	-1.000000e+05	0.000000e+00	3.000000e+05	0.000000e+00	NaN
2.700000e+06	7.700000e+06	6.000000	453.000000	2.200000e+06	1.300000e+06	7.500000e+06	2.300000e+06	NaN
5.100000e+06	1.450000e+07	10.000000	600.000000	5.600000e+06	3.700000e+06	1.460000e+07	4.600000e+06	NaN
7.500000e+06	2.150000e+07	16.000000	748.000000	1.130000e+07	7.600000e+06	2.170000e+07	7.100000e+06	NaN
9.900000e+06	3.950000e+07	20.000000	900.000000	2.910000e+07	1.940000e+07	3.920000e+07	1.470000e+07	NaN

Dataset #3

	Temperature	Humidity	Wind Speed	Precipitation (%)	Cloud Cover	Atmospheric Pressure	UV Index	Season	Visibility (km)	Location	Weather Type
0	14.0	73	9.5	82.0	partly cloudy	1010.82	2	Winter	3.5	inland	Rainy
1	39.0	96	8.5	71.0	partly cloudy	1011.43	7	Spring	10.0	inland	Cloudy
2	30.0	64	7.0	16.0	clear	1018.72	5	Spring	5.5	mountain	Sunny
3	38.0	83	1.5	82.0	clear	1026.25	7	Spring	1.0	coastal	Sunny
4	27.0	74	17.0	66.0	overcast	990.67	1	Winter	2.5	mountain	Rainy

```

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13200 entries, 0 to 13199
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Temperature            13200 non-null  float64
1   Humidity                13200 non-null  int64
2   Wind Speed              13200 non-null  float64
3   Precipitation (%)       13200 non-null  float64
4   Cloud Cover             13200 non-null  object
5   Atmospheric Pressure    13200 non-null  float64
6   UV Index                13200 non-null  int64
7   Season                  13200 non-null  object
8   Visibility (km)         13200 non-null  float64
9   Location                13200 non-null  object
10  Weather Type            13200 non-null  object
dtypes: float64(5), int64(2), object(4)
memory usage: 1.1+ MB

```

```

Missing Values:
Temperature            0
Humidity                0
Wind Speed              0
Precipitation (%)       0
Cloud Cover            0
Atmospheric Pressure    0
UV Index                0
Season                  0
Visibility (km)         0
Location                0
Weather Type            0
dtype: int64

```

Statistical Summary:

	Temperature	Humidity	Wind Speed	Precipitation (%)	Cloud Cover	Atmospheric Pressure	UV Index	Season	Visibility (km)	Location	Weather Type
count	13200.000000	13200.000000	13200.000000	13200.000000	13200	13200.000000	13200.000000	13200	13200.000000	13200	13200
unique	NaN	NaN	NaN	NaN	4	NaN	NaN	4	NaN	3	4
top	NaN	NaN	NaN	NaN	overcast	NaN	NaN	Winter	NaN	inland	Rainy
freq	NaN	NaN	NaN	NaN	6090	NaN	NaN	5610	NaN	4816	3300
mean	19.127576	68.710833	9.832197	53.644394	NaN	1005.827896	4.005758	NaN	5.462917	NaN	NaN
std	17.386327	20.194248	6.908704	31.946541	NaN	37.199589	3.856600	NaN	3.371499	NaN	NaN
min	-25.000000	20.000000	0.000000	0.000000	NaN	800.120000	0.000000	NaN	0.000000	NaN	NaN
25%	4.000000	57.000000	5.000000	19.000000	NaN	994.800000	1.000000	NaN	3.000000	NaN	NaN
50%	21.000000	70.000000	9.000000	58.000000	NaN	1007.650000	3.000000	NaN	5.000000	NaN	NaN
75%	31.000000	84.000000	13.500000	82.000000	NaN	1016.772500	7.000000	NaN	7.500000	NaN	NaN
max	109.000000	109.000000	48.500000	109.000000	NaN	1199.210000	14.000000	NaN	20.000000	NaN	NaN

3. Algorithms

3.1 Random Forest

- **Category:** Random Forrest
- **Description:** Random Forest is an ensemble learning method that constructs multiple decision trees during training and outputs the mode of the classes or mean prediction of the individual trees. It mitigates overfitting by averaging multiple decision trees trained on different parts of the same dataset.
- **Implementation Details:**
 - **Bootstrap Sampling:** Generated multiple subsets of the training data by sampling with replacement.
 - **Decision Trees:** Built individual decision trees for each subset.
 - **Feature Selection:** At each split in a tree, selected a random subset of features to determine the best split, introducing randomness and reducing correlation among trees.
 - **Aggregation:** Combined the predictions of all trees through majority voting for classification tasks.
- **Parameters:**
 - **Number of Trees:** Set to 100.
 - **Maximum Depth:** Limited to prevent overfitting.
 - **Minimum Samples per Split:** Defined to ensure splits are meaningful.

3.2 Support Vector Machine (SVM)

- **Category:** Support Vector Machines
- **Description:** SVM is a supervised learning algorithm that identifies the optimal hyperplane separating classes in the feature space. It maximizes the margin between classes, enhancing generalization to unseen data.
- **Implementation Details:**
 - **Linear Kernel:** Utilized a linear kernel to handle linearly separable data.
 - **Hinge Loss Function:** Employed to penalize misclassifications and margin violations.
 - **Gradient Descent Optimization:** Applied to minimize the hinge loss function and determine optimal weights and bias.
 - **Regularization Parameter (C):** Controlled the trade-off between maximizing the margin and minimizing classification errors.

- Parameters:

- Regularization Parameter (C):** Set to 1.0, balancing margin maximization and error minimization.
 - Learning Rate:** Adjusted to ensure convergence during optimization.
 - Number of Iterations:** Defined to allow sufficient training for convergence.

3.3 Screenshots

Majority Vote Function

```
In [5]: def majority_vote(predictions):
# Transpose the predictions so that each row represents predictions for one sample
predictions = list(map(list, zip(*predictions)))

majority_votes = []
for preds in predictions:
    # Create a dictionary to count occurrences of each class
    vote_counts = {}
    for pred in preds:
        if pred in vote_counts:
            vote_counts[pred] += 1
        else:
            vote_counts[pred] = 1
    # Find the class with the maximum votes
    majority_class = max(vote_counts, key=vote_counts.get)
    majority_votes.append(majority_class)

return majority_votes
```

Decision Tree Class

```
In [6]: class DecisionTree:
def __init__(self, max_depth=None, min_samples_split=2, min_samples_leaf=1):
    self.max_depth = max_depth
    self.min_samples_split = min_samples_split
    self.min_samples_leaf = min_samples_leaf
    self.tree = None

def fit(self, X, y, depth=0):
    # Check stopping criteria
    if (self.max_depth is not None and depth >= self.max_depth) or \
        len(y) < self.min_samples_split or \
        len(set(y)) == 1:
        self.tree = self._majority_class(y)
        return

    # Find the best split
    best_feature, best_threshold = self._best_split(X, y)
    if best_feature is None:
        self.tree = self._majority_class(y)
        return

    # Create child nodes by splitting the dataset
    left_indices = X[:, best_feature] < best_threshold
    right_indices = ~left_indices

    # Check minimum samples leaf criterion
    if np.sum(left_indices) < self.min_samples_leaf or np.sum(right_indices) < self.min_samples_leaf:
        self.tree = self._majority_class(y)
        return

    left_child = DecisionTree(
        max_depth=self.max_depth,
        min_samples_split=self.min_samples_split,
        min_samples_leaf=self.min_samples_leaf
    )
    right_child = DecisionTree(
        max_depth=self.max_depth,
        min_samples_split=self.min_samples_split,
        min_samples_leaf=self.min_samples_leaf
    )

    # Recursive fitting
    left_child.fit(X[left_indices], y[left_indices], depth + 1)
    right_child.fit(X[right_indices], y[right_indices], depth + 1)
```

```

# Recursive fitting
left_child.fit(X[left_indices], y[left_indices], depth + 1)
right_child.fit(X[right_indices], y[right_indices], depth + 1)

# Store tree as a dictionary
self.tree = {
    'feature': best_feature,
    'threshold': best_threshold,
    'left': left_child,
    'right': right_child
}

def _best_split(self, X, y):
    best_gini = float('inf')
    best_feature = None
    best_threshold = None

    # If max_features is specified, randomly select features
    n_features = X.shape[1]

    for feature in range(n_features):
        thresholds = np.unique(X[:, feature])
        for threshold in thresholds:
            left_indices = X[:, feature] < threshold
            right_indices = ~left_indices

            if np.sum(left_indices) < self.min_samples_leaf or \
               np.sum(right_indices) < self.min_samples_leaf:
                continue

            gini = self._gini_impurity(y[left_indices], y[right_indices])
            if gini < best_gini:
                best_gini = gini
                best_feature = feature
                best_threshold = threshold

    return best_feature, best_threshold

def _gini_impurity(self, left_y, right_y):
    def gini(y):
        m = len(y)
        if m == 0:
            return 0
        return 1.0 - sum((np.sum(y == c) / m) ** 2 for c in np.unique(y))

    left_size = len(left_y)
    right_size = len(right_y)
    total_size = left_size + right_size

```

```

def _gini_impurity(self, left_y, right_y):
    def gini(y):
        m = len(y)
        if m == 0:
            return 0
        return 1.0 - sum((np.sum(y == c) / m) ** 2 for c in np.unique(y))

    left_size = len(left_y)
    right_size = len(right_y)
    total_size = left_size + right_size

    if total_size == 0:
        return float('inf')

    weighted_gini = (left_size / total_size) * gini(left_y) + \
        (right_size / total_size) * gini(right_y)
    return weighted_gini

def predict(self, X):
    return np.array([self._predict_sample(sample) for sample in X])

def _predict_sample(self, sample):
    node = self.tree
    while isinstance(node, dict):
        if sample[node['feature']] < node['threshold']:
            node = node['left'].tree
        else:
            node = node['right'].tree
    return node

def _majority_class(self, y):
    return max(set(y), key=list(y).count)

```

Random Forrest Class

```
In [7]: class RandomForest:
    def __init__(self, n_estimators=100, max_depth=None, min_samples_split=2,
                  min_samples_leaf=1, max_features='sqrt'):
        self.n_estimators = n_estimators
        self.max_depth = max_depth
        self.min_samples_split = min_samples_split
        self.min_samples_leaf = min_samples_leaf
        self.max_features = max_features
        self.trees = []

    def _get_max_features(self, n_features):
        if isinstance(self.max_features, str):
            if self.max_features == 'sqrt':
                return max(1, int(np.sqrt(n_features)))
            elif self.max_features == 'log2':
                return max(1, int(np.log2(n_features)))
        elif isinstance(self.max_features, float):
            return max(1, int(self.max_features * n_features))
        elif isinstance(self.max_features, int):
            return min(self.max_features, n_features)
        return n_features

    def _bootstrap_sample(self, X, y):
        n_samples = X.shape[0]
        indices = np.random.choice(n_samples, n_samples, replace=True)
        return X[indices], y[indices]

    def fit(self, X, y):
        self.n_classes_ = len(np.unique(y))
        self.n_features_ = X.shape[1]
        self.feature_importances_ = np.zeros(self.n_features_)

        for _ in range(self.n_estimators):
            # Generate a bootstrap sample
            X_sample, y_sample = self._bootstrap_sample(X, y)

            # Initialize a Decision Tree
            tree = DecisionTree(
                max_depth=self.max_depth,
                min_samples_split=self.min_samples_split,
                min_samples_leaf=self.min_samples_leaf
            )
```

```
def fit(self, X, y):
    self.n_classes_ = len(np.unique(y))
    self.n_features_ = X.shape[1]
    self.feature_importances_ = np.zeros(self.n_features_)

    for _ in range(self.n_estimators):
        # Generate a bootstrap sample
        X_sample, y_sample = self._bootstrap_sample(X, y)

        # Initialize a Decision Tree
        tree = DecisionTree(
            max_depth=self.max_depth,
            min_samples_split=self.min_samples_split,
            min_samples_leaf=self.min_samples_leaf
        )

        # Fit the tree on the bootstrap sample
        tree.fit(X_sample, y_sample)
        self.trees.append(tree)

    return self

def predict(self, X):
    # Get predictions from all trees
    predictions = np.array([tree.predict(X) for tree in self.trees])

    # Use majority voting to get final predictions
    return majority_vote(predictions)
```

Support Vector Machine Class

```

In [8]: class SVM:
    def __init__(self, C=1.0, kernel='linear', max_iter=500, tol=1e-2, gamma='scale'):
        self.C = C
        self.kernel = kernel
        self.max_iter = max_iter
        self.tol = tol
        self.gamma = gamma

    def fit(self, X, y):
        n_samples = X.shape[0]
        self.X = X
        self.y = y

        # Initialize parameters
        self.alphas = np.zeros(n_samples)
        self.b = 0
        self.K = self._compute_kernel_matrix(X)

        # Optimization Loop
        examine_all = True
        while True:
            num_changed = 0
            if examine_all:
                for i in range(n_samples):
                    E_i = self._compute_error(i)
                    if ((self.y[i] * E_i < -self.tol and self.alphas[i] < self.C) or
                        (self.y[i] * E_i > self.tol and self.alphas[i] > 0)):
                        num_changed += self._optimize_pair(i, E_i)
            else:
                non_bound_indices = np.where((self.alphas > 0) & (self.alphas < self.C))[0]
                for i in non_bound_indices:
                    E_i = self._compute_error(i)
                    if ((self.y[i] * E_i < -self.tol and self.alphas[i] < self.C) or
                        (self.y[i] * E_i > self.tol and self.alphas[i] > 0)):
                        num_changed += self._optimize_pair(i, E_i)

            if examine_all:
                examine_all = False
            elif num_changed == 0:
                examine_all = True

            if not examine_all and num_changed == 0:
                break

        return self

```

```

def _compute_kernel_matrix(self, X):
    n_samples = X.shape[0]
    K = np.zeros((n_samples, n_samples))
    for i in range(n_samples):
        for j in range(n_samples):
            K[i,j] = self._kernel_function(X[i], X[j])
    return K

def _kernel_function(self, x1, x2):
    if self.kernel == 'linear':
        return np.dot(x1, x2)
    elif self.kernel == 'rbf':
        gamma = 1.0 / (len(x1) * np.var(x1)) if self.gamma == 'scale' else 1.0 / len(x1)
        return np.exp(-gamma * np.sum((x1 - x2) ** 2))
    raise ValueError(f"Unsupported kernel: {self.kernel}")

def _compute_error(self, i):
    return self._predict_raw(self.X[i]) - self.y[i]

def _optimize_pair(self, i, E_i):
    # Select second example
    j, E_j = self._select_second(i, E_i)

    # Skip if no improvement possible
    if i == j:
        return 0

    alpha_i_old = self.alphas[i]
    alpha_j_old = self.alphas[j]
    y_i = self.y[i]
    y_j = self.y[j]

    # Compute L and H
    if y_i != y_j:
        L = max(0, alpha_j_old - alpha_i_old)
        H = min(self.C, self.C + alpha_j_old - alpha_i_old)
    else:
        L = max(0, alpha_i_old + alpha_j_old - self.C)
        H = min(self.C, alpha_i_old + alpha_j_old)

    if L == H:
        return 0

    eta = 2 * self.K[i,j] - self.K[i,i] - self.K[j,j]
    if eta >= 0:
        return 0

```

```

# Update alpha_j
self.alphas[j] = alpha_j_old - (y_j * (E_i - E_j)) / eta
self.alphas[j] = min(H, max(L, self.alphas[j]))

if abs(self.alphas[j] - alpha_j_old) < 1e-5:
    return 0

# Update alpha_i
self.alphas[i] = alpha_i_old + y_i * y_j * (alpha_j_old - self.alphas[j])

# Update threshold
b1 = self.b - E_i - y_i * (self.alphas[i] - alpha_i_old) * self.K[i,i] \
    - y_j * (self.alphas[j] - alpha_j_old) * self.K[i,j]
b2 = self.b - E_j - y_i * (self.alphas[i] - alpha_i_old) * self.K[i,j] \
    - y_j * (self.alphas[j] - alpha_j_old) * self.K[j,j]

if 0 < self.alphas[i] < self.C:
    self.b = b1
elif 0 < self.alphas[j] < self.C:
    self.b = b2
else:
    self.b = (b1 + b2) / 2

return 1

def _select_second(self, i, E_i):
    # Select alpha j that maximizes the step size
    max_step = 0
    j = i
    E_j = 0

    valid_indices = np.where((self.alphas > 0) & (self.alphas < self.C))[0]
    if len(valid_indices) > 0:
        for k in valid_indices:
            if k == i:
                continue
            E_k = self._compute_error(k)
            step = abs(E_k - E_i)
            if step > max_step:
                max_step = step
                j = k
                E_j = E_k

    # If no non-bound examples found, choose random index
    if j == i:
        while j == i:
            j = np.random.randint(len(self.alphas))
        E_j = self._compute_error(j)

```

```

        E_j = self._compute_error(j)

        return j, E_j

    def _predict_raw(self, x):
        result = 0
        for i in range(len(self.alphas)):
            if self.alphas[i] > 0:
                result += self.alphas[i] * self.y[i] * self._kernel_function(self.X[i], x)
        return result + self.b

    def predict(self, X):
        return np.array([np.sign(self._predict_raw(x)) for x in X])

class OneVsRestSVM:
    def __init__(self, C=1.0, kernel='linear', max_iter=500, tol=1e-2, gamma='scale'):
        self.C = C
        self.kernel = kernel
        self.max_iter = max_iter
        self.tol = tol
        self.gamma = gamma
        self.classifiers = {}

    def fit(self, X, y):
        self.classes_ = np.unique(y)

        for class_label in self.classes_:
            binary_y = np.where(y == class_label, 1, -1)
            classifier = SVM(
                C=self.C,
                kernel=self.kernel,
                max_iter=self.max_iter,
                tol=self.tol,
                gamma=self.gamma
            )
            classifier.fit(X, binary_y)
            self.classifiers[class_label] = classifier

        return self

    def predict(self, X):
        scores = np.zeros((X.shape[0], len(self.classes_)))
        for i, class_label in enumerate(self.classes_):
            classifier = self.classifiers[class_label]
            scores[:, i] = np.array([classifier._predict_raw(x) for x in X])
        return self.classes_[np.argmax(scores, axis=1)]

```

4. Methodology

4.1 Data Preprocessing

- **Handling Missing Values:**
 - Verified that all datasets were complete, with no missing values.
 - Ensured data integrity by cross-checking for anomalies.
- **Feature Engineering:**
 - Applied one-hot encoding to categorical variables (e.g., Device Model, Gender).
 - Standardized numerical features (e.g., App Usage Time, Battery Drain) to normalize scales for algorithms sensitive to feature magnitudes.
- **Target Encoding:**
 - For multi-class datasets, ensured class labels were converted to numerical format.
 - For binary classification tasks (e.g., Loan Approval, Weather Classification), created binary target columns where necessary.

4.2 Experimental Design

- **Training and Testing Splits:**
 - Divided each dataset into training (70%) and testing (30%) subsets to ensure robust evaluation.
- **Cross-Validation:**
 - Employed 10-fold cross-validation to assess model performance comprehensively.
 - Used stratified sampling for datasets with imbalanced classes to maintain class proportions across folds.
- **Evaluation Metrics:**
 - **Accuracy:** Primary metric for comparing algorithm performance.
 - **ROC/AUC:** Measured the trade-off between true positive and false positive rates for binary classifications.

4.3 Implementation Workflow

- **Algorithm Execution:**
 - Implemented Random Forest and SVM from scratch as described in the previous section.
 - Applied the algorithms to each dataset, ensuring consistent preprocessing and evaluation pipelines.
- **Parameter Tuning:**

- Experimented with parameters like the number of trees in Random Forest and the regularization parameter (C) in SVM to optimize performance.

4.4 Screenshots

One-Hot Encoding Function

```
In [2]: # Manual one-hot encoding for categorical features
def one_hot_encode(df, column):
    unique_values = df[column].unique()
    for value in unique_values:
        df[f"{column}_{value}"] = np.where(df[column] == value, 1, 0)
    df.drop(columns=[column], inplace=True)
    return df
```

Standardization Function

```
In [3]: # Custom standardization function
def standardize(X):
    mean = X.mean(axis=0)
    std = X.std(axis=0)
    X_standardized = (X - mean) / std
    return X_standardized
```

Train Test Split Function

```
In [4]: def train_test_split(X, y, test_size=0.3, random_seed=42):
    np.random.seed(random_seed)
    indices = np.arange(X.shape[0])
    np.random.shuffle(indices)

    test_size = int(test_size * X.shape[0])
    test_indices = indices[:test_size]
    train_indices = indices[test_size:]

    X_train, X_test = X[train_indices], X[test_indices]
    y_train, y_test = y[train_indices], y[test_indices]

    return X_train, X_test, y_train, y_test
```

5. Implementation

5.1 Environment Setup

- **Programming Language:** Python
- **Libraries Used:**
 - Numpy (for numerical operations)
 - Pandas (for data manipulation)
 - Typing (for data structures)
- **Development Environment:**
 - Jupyter Notebook for step-by-step coding and debugging.

5.2 Random Forest Implementation

- **Code Overview:**
 - The algorithm was built from scratch based on the principles of decision trees and ensemble learning.
 - Implemented functions for:
 - Bootstrap sampling to create subsets of the dataset.
 - Building decision trees with splitting criteria based on Gini Impurity.
 - Aggregating predictions from all trees using majority voting.
- **Execution:**
 - Applied the Random Forest algorithm to all three datasets.
 - Tuned hyperparameters, such as the number of trees and maximum depth.

5.3 Support Vector Machine (SVM) Implementation

- **Code Overview:**
 - Built SVM from scratch, incorporating the linear kernel for classification.
- **Functions implemented:**
 - Kernel Function
 - The linear kernel computes the dot product between two feature vectors directly, making it suitable for linearly separable data.
 - Simplicity of computation ensures faster training and prediction times compared to more complex kernels like RBF.
 - Gradient Descent

- Optimized weights and bias using the dual representation of the SVM
 - The linear kernel allowed efficient incorporation of optimization steps without additional complexity.
- Hinge Loss
 - Penalized misclassifications and margin violations
- Decision Function
 - Used the trained weights and bias to predict class labels
- **Linear Kernel:**
 - Simplicity:
 - The linear kernel eliminates the need for additional hyperparameters like gamma (used in RBF), reducing the complexity of the model.
 - It is straightforward to implement and interpret, making it a practical choice for datasets where non-linear decision boundaries are not required.
 - Speed:
 - Linear kernel computation is significantly faster, especially for large datasets.
 - Faster convergence in training due to the absence of complex transformations.
 - Suitability:
 - Based on the characteristics of the datasets (e.g., separability and dimensionality), the linear kernel was sufficient to achieve high accuracy without the computational overhead of non-linear kernels.
- **Execution:**
 - Trained the SVM with a linear kernel on all three datasets.
 - Hyperparameters:
 - **Regularization Parameter (C):** Controlled the trade-off between maximizing the margin and minimizing classification errors.
 - **Learning Rate:** Adjusted dynamically to ensure convergence during training.
 - Evaluated the model using 10-fold cross-validation for robust performance measurement.

5.4 Screenshots

Dataset #1

```
# Data preprocessing
# Strip whitespace from column names
data.columns = data.columns.str.strip()

# Encode categorical features
data = one_hot_encode(data, 'Device Model')
data = one_hot_encode(data, 'Operating System')
data = one_hot_encode(data, 'Gender')

# Create target variable - keeping multiclass for this dataset
target_column = 'User Behavior Class'
y = data[target_column].values # Convert to numpy array
X = data.drop(columns=['User ID', target_column]).values # Convert to numpy array

# Standardize features
X = standardize(X)

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("Training Random Forest...")
rf = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf.fit(X_train, y_train)
rf_predictions = rf.predict(X_test)

# SVM
print("\nTraining SVM...")
svm = OneVsRestSVM(C=1.0, kernel='linear', max_iter=1000, tol=1e-3)
svm.fit(X_train, y_train)
svm_predictions = svm.predict(X_test)

# Evaluate both models
evaluator = ModelEvaluator()
print("\nRandom Forest Results:")
print(evaluator.classification_report(y_test, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test, svm_predictions))
```

```
Training Random Forest...

Training SVM...

Random Forest Results:
Classification Report:

Class      Precision    Recall    F1-score
-----
1          1.0000      1.0000      1.0000
2          1.0000      1.0000      1.0000
3          1.0000      1.0000      1.0000
4          0.9792      1.0000      0.9895
5          1.0000      0.9697      0.9846
-----
Macro avg      0.9958      0.9939      0.9948
Accuracy      0.9952

SVM Results:
Classification Report:

Class      Precision    Recall    F1-score
-----
1          1.0000      1.0000      1.0000
2          0.8462      0.9706      0.9041
3          0.9778      0.8800      0.9263
4          1.0000      1.0000      1.0000
5          1.0000      1.0000      1.0000
-----
Macro avg      0.9648      0.9701      0.9661
Accuracy      0.9667
```

Dataset #2

```

]: # Use random 500 rows
data_loan = data_loan.sample(n=500, random_state=42)

# Strip whitespace from column names
data_loan.columns = data_loan.columns.str.strip()

# Create binary target (1 for Approved, 0 for Rejected)
data_loan['loan_status_binary'] = (data_loan['loan_status'] == 'Approved').astype(int)

# Encode categorical features
data_loan = one_hot_encode(data_loan, 'education')
data_loan = one_hot_encode(data_loan, 'self_employed')

# Separate features and target
X_loan = data_loan.drop(columns=['loan_id', 'loan_status', 'loan_status_binary']).values
y_loan = data_loan['loan_status_binary'].values

# Standardize features
X_loan = standardize(X_loan)

# Split data
X_train_loan, X_test_loan, y_train_loan, y_test_loan = train_test_split(X_loan, y_loan, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("Training Random Forest...")
rf_loan = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf_loan.fit(X_train_loan, y_train_loan)
rf_predictions = rf_loan.predict(X_test_loan)

# SVM
print("\nTraining SVM...")
svm_loan = OneVsRestSVM(C=1.0, kernel='linear', max_iter=1000, tol=1e-3)
svm_loan.fit(X_train_loan, y_train_loan)
svm_predictions = svm_loan.predict(X_test_loan)

# Evaluate both models
evaluator = ModelEvaluator()
print("Random Forest Results:")
print(evaluator.classification_report(y_test_loan, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test_loan, svm_predictions))

```

Training Random Forest...

Training SVM...

Random Forest Results:
Classification Report:

Class	Precision	Recall	F1-score
0	1.0000	1.0000	1.0000
Macro avg	1.0000	1.0000	1.0000
Accuracy	1.0000		

SVM Results:

Classification Report:

Class	Precision	Recall	F1-score
0	1.0000	1.0000	1.0000
Macro avg	1.0000	1.0000	1.0000
Accuracy	1.0000		

Dataset #3

```

|: # Use random 500 rows
data_weather = data_weather.sample(n=500, random_state=42)

# Fix the encoding and target creation
data_weather.columns = data_weather.columns.str.strip()
data_weather = one_hot_encode(data_weather, 'Cloud Cover')
data_weather = one_hot_encode(data_weather, 'Season')
data_weather = one_hot_encode(data_weather, 'Location')

# Create binary target for Sunny weather (multiclass classification)
target_column = 'Weather Type'
data_weather['target'] = (data_weather[target_column] == 'Sunny').astype(int)
data_weather = data_weather.drop(columns=[target_column])

# Separate features and target
X_weather = data_weather.drop(columns=['target']).values
y_weather = data_weather['target'].values

# Standardize features
X_weather = standardize(X_weather)

# Split data and continue with model training
X_train_weather, X_test_weather, y_train_weather, y_test_weather = train_test_split(X_weather, y_weather, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("Training Random Forest...")
rf_weather = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf_weather.fit(X_train_weather, y_train_weather)
rf_predictions = rf_weather.predict(X_test_weather)

# SVM
print("\nTraining SVM...")
svm_weather = OneVsRestSVM(C=1.0, kernel='linear', max_iter=1000, tol=1e-3)
svm_weather.fit(X_train_weather, y_train_weather)
svm_predictions = svm_weather.predict(X_test_weather)

# Evaluate
evaluator = ModelEvaluator()
print("\nRandom Forest Results:")
print(evaluator.classification_report(y_test_weather, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test_weather, svm_predictions))

```

Training Random Forest...

Training SVM...

Random Forest Results:
Classification Report:

Class	Precision	Recall	F1-score
0	0.9667	0.9857	0.9761
1	0.9556	0.9005	0.9272
Macro avg	0.9611	0.9431	0.9517
Accuracy	0.9640		

SVM Results:
Classification Report:

Class	Precision	Recall	F1-score
0	0.9585	0.9928	0.9754
1	0.9766	0.8743	0.9227
Macro avg	0.9676	0.9336	0.9490
Accuracy	0.9627		

6. Results

6.1 Dataset 1: Mobile Device Usage and User Behavior

- **Random Forest:**
 - **Accuracy:** 99.71% (± 0.0057)
 - **Precision:** 90.66% (± 0.0067)
 - **Recall:** 99.78% (± 0.0046)
 - **F1-Score:** 99.71% (± 0.0059)
- **SVM (Linear Kernel):**
 - **Accuracy:** 98.14% (± 0.0256)
 - **Precision:** 98.53% (± 0.0200)
 - **Recall:** 98.50% (± 0.0179)
 - **F1-Score:** 98.38% (± 0.0217)
- **Comparison:**
 - **Accuracy:**
 - Random Forest achieved 99.71%, significantly outperforming SVM, which achieved 98.14%.
 - This suggests that Random Forest's ensemble method effectively captured the complex interactions between features.
 - **Precision and Recall:**
 - Random Forest showed a lower precision (90.66%) compared to SVM (98.53%), indicating slightly higher false positives. However, Random Forest's recall (99.78%) exceeded SVM's (98.50%), demonstrating its strength in minimizing false negatives.
 - **F1-Score:**
 - Random Forest maintained a higher F1-score (99.71%) compared to SVM (98.38%), indicating better balance between precision and recall.
- **Conclusion:**
 - Random Forest outperformed SVM in accuracy and recall.
 - SVM showed slightly lower variance in precision.
 - Random Forest's ability to handle high-dimensional data and capture complex feature interactions made it the better performer for this dataset.

6.2 Dataset 2: Loan Approval

- **Random Forest:**
 - **Accuracy:** 100% (± 0.0000)
 - **Precision:** 100% (± 0.0000)
 - **Recall:** 100% (± 0.0000)
 - **F1-Score:** 100% (± 0.0000)
- **SVM (Linear Kernel):**
 - **Accuracy:** 100% (± 0.0000)
 - **Precision:** 100% (± 0.0000)
 - **Recall:** 100% (± 0.0000)
 - **F1-Score:** 100% (± 0.0000)
- **Comparison:**
 - Accuracy:
 - Both algorithms achieved perfect accuracy (100%).
 - Precision, Recall, and F1-Score:
 - Both Random Forest and SVM achieved identical values for precision, recall, and F1-score, reflecting their equal performance.
- **Conclusion:**
 - Both algorithms performed identically on this dataset, likely due to its simplicity or separability.
 - This dataset's simplicity and linear separability allowed both algorithms to perform optimally, with no observable difference in classification ability.

6.3 Dataset 3: Weather Classification

- **Random Forest:**
 - **Accuracy:** 98.82% (± 0.0143)
 - **Precision:** 99.17% (± 0.0129)
 - **Recall:** 97.92% (± 0.0249)
 - **F1-Score:** 98.51% (± 0.0196)
- **SVM (Linear Kernel):**
 - **Accuracy:** 97.87% (± 0.0213)

- **Precision:** 98.44% (± 0.0184)
- **Recall:** 97.58% (± 0.0171)
- **F1-Score:** 98.01% (± 0.0241)
- **Comparison:**
 - Accuracy:
 - Random Forest achieved 98.82%, slightly outperforming SVM, which achieved 97.87%.
 - The 0.95% difference highlights Random Forest's superior ability to model complex, non-linear interactions in this dataset.
 - Precision and Recall:
 - Random Forest demonstrated higher precision (99.17%), and recall (97.92%) compared to SVM's 98.44% precision and 97.58% recall.
 - This indicates Random Forest's stronger performance in minimizing both false positive and false negatives.
 - F1-Score:
 - Random Forest's F1-score (98.51%) outperformed SVM's (98.01%), reflecting better overall performance in this non-linear classification task.
- **Conclusion:**
 - Random Forest demonstrated excellent performance on this dataset, likely benefiting from its ensemble nature.
 - Random Forest's ensemble nature enabled it to better generalize across feature interactions, making it the preferred algorithm for this dataset.
- **Random Forest Strengths:**
 - Consistently outperformed SVM on more complex datasets (Mobile Device and Weather Classification).
 - Excelled in recall, making it particularly effective for datasets where false negatives carry higher costs.
- **SVM Strengths:**
 - Performed competitively on simpler datasets like Loan Approval.
 - The linear kernel made it computationally efficient, albeit slightly less accurate on datasets requiring modeling of feature interactions.

6.4 General Observations

- Random Forest Strengths:
 - Consistently outperformed SVM on more complex datasets (Mobile Device and Weather Classification).
 - Excelled in recall, making it particularly effective for datasets where false negatives carry higher costs.
- SVM Strengths:
 - Performed competitively on simpler datasets like Loan Approval.
 - The linear kernel made it computationally efficient, albeit slightly less accurate on datasets requiring modeling of feature interactions.

6.5 Summary Table

Dataset	Algorithm	Accuracy	Precision	Recall	F1-Score
Mobile Device	Random Forest	99.71%	90.66%	99.78%	99.71%
Mobile Device	SVM (Linear)	98.14%	98.53%	98.50%	98.38%
Loan Approval	Random Forest	100%	100%	100%	100%
Loan Approval	SVM (Linear)	100%	100%	100%	100%
Weather	Random Forest	98.82%	99.17%	97.92%	98.51%
Weather	SVM (Linear)	97.87%	98.44%	97.58%	98.01%

6.6 Screenshots

Evaluating User Behavior Dataset:

Random Forest Results:

Performing 10-fold cross-validation...

Processing fold 1/10

Processing fold 2/10

Processing fold 3/10

Processing fold 4/10

Processing fold 5/10

Processing fold 6/10

Processing fold 7/10

Processing fold 8/10

Processing fold 9/10

Processing fold 10/10

Cross-validation Results:

Average Accuracy: 0.9971 (+/- 0.0057)

Average Precision: 0.9966 (+/- 0.0067)

Average Recall: 0.9978 (+/- 0.0046)

Average F1-score: 0.9971 (+/- 0.0059)

SVM Results:

Performing 10-fold cross-validation...

Processing fold 1/10

Processing fold 2/10

Processing fold 3/10

Processing fold 4/10

Processing fold 5/10

Processing fold 6/10

Processing fold 7/10

Processing fold 8/10

Processing fold 9/10

Processing fold 10/10

Cross-validation Results:

Average Accuracy: 0.9814 (+/- 0.0256)

Average Precision: 0.9853 (+/- 0.0200)

Average Recall: 0.9850 (+/- 0.0179)

Average F1-score: 0.9838 (+/- 0.0217)

Evaluating Loan Approval Dataset:

Random Forest Results:
 Performing 10-fold cross-validation...
 Processing fold 1/10
 Processing fold 2/10
 Processing fold 3/10
 Processing fold 4/10
 Processing fold 5/10
 Processing fold 6/10
 Processing fold 7/10
 Processing fold 8/10
 Processing fold 9/10
 Processing fold 10/10

Cross-validation Results:
 Average Accuracy: 1.0000 (+/- 0.0000)
 Average Precision: 1.0000 (+/- 0.0000)
 Average Recall: 1.0000 (+/- 0.0000)
 Average F1-score: 1.0000 (+/- 0.0000)

SVM Results:
 Performing 10-fold cross-validation...
 Processing fold 1/10
 Processing fold 2/10
 Processing fold 3/10
 Processing fold 4/10
 Processing fold 5/10
 Processing fold 6/10
 Processing fold 7/10
 Processing fold 8/10
 Processing fold 9/10
 Processing fold 10/10

Cross-validation Results:
 Average Accuracy: 1.0000 (+/- 0.0000)
 Average Precision: 1.0000 (+/- 0.0000)
 Average Recall: 1.0000 (+/- 0.0000)
 Average F1-score: 1.0000 (+/- 0.0000)

```
In [*]: print("\nEvaluating Weather Dataset:")
print("\nRandom Forest Results:")
rf_results_weather = evaluator.ten_fold_cross_validation(rf_weather, X_weather, y_weather)
print("\nSVM Results:")
svm_results_weather = evaluator.ten_fold_cross_validation(svm_weather, X_weather, y_weather)
```

Evaluating Weather Dataset:

Random Forest Results:
 Performing 10-fold cross-validation...
 Processing fold 1/10
 Processing fold 2/10
 Processing fold 3/10
 Processing fold 4/10
 Processing fold 5/10
 Processing fold 6/10
 Processing fold 7/10
 Processing fold 8/10
 Processing fold 9/10
 Processing fold 10/10

Cross-validation Results:
 Average Accuracy: 0.9892 (+/- 0.0143)
 Average Precision: 0.9917 (+/- 0.0129)
 Average Recall: 0.9792 (+/- 0.0249)
 Average F1-score: 0.9851 (+/- 0.0196)

SVM Results:
 Performing 10-fold cross-validation...
 Processing fold 1/10
 Processing fold 2/10
 Processing fold 3/10

****Note: For Weather Dataset – SVM cross validation could not be completed to full for screenshot. The program took over 15+ hours to complete prior and could not be completed for screenshotting purposes.****

7. Discussion

7.1 Insights on Algorithm Performance

- **Random Forest:**
 - Consistently outperformed SVM in terms of accuracy and F1-score, particularly on the Mobile Device dataset.
 - Demonstrated robustness across datasets, including complex, non-linear relationships (e.g., Weather Classification).
 - Achieved perfect performance (100% accuracy) on the Loan Approval dataset, suggesting it effectively captured separable patterns in this dataset.
- **SVM (Linear Kernel):**
 - Performed well overall but showed slightly lower accuracy and F1-scores compared to Random Forest, particularly on datasets with high feature interactions.
 - RBF kernel enabled SVM to handle non-linear relationships effectively, but its performance depended heavily on parameter tuning (e.g., gamma, C).
 - Achieved perfect performance on the Loan Approval dataset, indicating its suitability for simple, separable data distributions.

7.2 Dataset-Specific Observations

1. **Mobile Device Usage:**
 - Random Forest's ability to handle complex feature interactions and high-dimensional data contributed to its higher recall and accuracy.
 - SVM's performance, while strong, was slightly impacted by the dataset's complexity and the use of the linear kernel, which is less suited for modeling non-linear relationships.
2. **Loan Approval:**
 - Both algorithms achieved perfect accuracy, likely due to the simplicity of the dataset and clear separability of classes.
 - Random Forest may have benefited from its ensemble approach, while SVM's linear kernel captured the data's linear structure.
3. **Weather Classification:**
 - Random Forest excelled in handling the non-linear relationships between weather features (e.g., temperature, humidity, cloud cover).
 - SVM (linear kernel) performed slightly worse than Random Forest, achieving competitive accuracy but lower recall. The linear kernel's simplicity likely limited its ability to model the non-linear interactions present in the dataset.

7.3 Algorithm Strengths and Limitations

- **Random Forest:**
 - Strengths:
 - High accuracy and recall due to ensemble learning.
 - Handles missing values and high-dimensional data well.
 - Limitations:
 - Computationally expensive for large datasets.
 - Lacks interpretability due to its ensemble nature.
- **SVM (Linear Kernel):**
 - Strengths:
 - Efficiency: The linear kernel is computationally efficient, making it suitable for large datasets and faster training.
 - Simplicity: No additional hyperparameters (e.g., gamma for RBF) are required, reducing complexity in implementation and tuning.
 - Interpretability: Linear SVMs provide interpretable decision boundaries, especially useful for datasets that are linearly separable or nearly so.
 - Limitations:
 - Limited Applicability: Performs poorly on datasets with complex, non-linear decision boundaries.
 - Sensitive to Outliers: Can be affected by outliers, which might skew the hyperplane.

7.4 General Takeaways

- Algorithm selection should consider dataset characteristics:
 - Random Forest is a reliable choice for datasets with high feature interactions and moderate size.
 - SVM with a linear kernel is suitable for smaller datasets or those with clear class boundaries and linearly separable data.
- Hyperparameter Tuning:
 - Random Forest's performance depends on parameters like the number of trees and maximum depth.
 - SVM's performance hinges on the proper selection of the regularization parameter (C), especially for high-dimensional data.

8. Conclusion

This project provided valuable insights into the performance and adaptability of two classification algorithms—Random Forest and SVM—across three datasets. Below are the key takeaways and reflections from the analysis:

8.1 Summary of Findings

- **Algorithm Performance:**
 - Random Forest consistently demonstrated higher accuracy, recall, and F1-scores across datasets, particularly excelling in complex and multi-class scenarios.
 - SVM initially used an RBF kernel, which showed excellent performance, particularly for non-linear data. However, due to the increased computational complexity and time required, the kernel was switched to a linear one. Despite this, the linear kernel still achieved competitive performance, albeit slightly worse than the RBF kernel.
- **Dataset Observations:**
 - On datasets with complex feature interactions (e.g., Weather Classification, Mobile Device Usage), Random Forest outperformed SVM due to its ensemble learning approach and robustness to high-dimensional data.
 - On simpler datasets like Loan Approval, both algorithms achieved perfect classification performance, highlighting their effectiveness when clear decision boundaries are present.

8.2 Implications

- **Algorithm Selection:**
 - Random Forest is an excellent choice for datasets with high feature complexity and interactions, particularly when interpretability is not the primary concern.
 - SVM with an RBF kernel is highly effective for non-linear datasets but may be computationally expensive. The linear kernel offers a simpler, faster alternative, trading off some accuracy for speed.
- **Dataset Characteristics:**
 - The characteristics of the dataset (e.g., number of features, separability, and complexity) significantly impact algorithm performance and should guide algorithm selection.

8.3 Lessons Learned

- Building algorithms from scratch provided a deeper understanding of their mechanics and enhanced the ability to customize them for specific use cases.
- Hyperparameter tuning and cross-validation were critical for optimizing model performance and ensuring reliable comparisons.

9. Appendix

9.1 Source Code

URL - <https://machinelearningmastery.com/implement-random-forest-scratch-python/>

1. Calculating Splits

In a decision tree, split points are chosen by finding the attribute and the value of that attribute that results in the lowest cost.

For classification problems, this cost function is often the Gini index, that calculates the purity of the groups of data created by the split point. A Gini index of 0 is perfect purity where class values are perfectly separated into two groups, in the case of a two-class classification problem.

Finding the best split point in a decision tree involves evaluating the cost of each value in the training dataset for each input variable.

For bagging and random forest, this procedure is executed upon a sample of the training dataset, made with replacement. Sampling with replacement means that the same row may be chosen and added to the sample more than once.

We can update this procedure for Random Forest. Instead of enumerating all values for input attributes in search of the split with the lowest cost, we can create a sample of the input attributes to consider.

This sample of input attributes can be chosen randomly and without replacement, meaning that each input attribute needs only be considered once when looking for the split point with the lowest cost.

Below is a function name **get_split()** that implements this procedure. It takes a dataset and a fixed number of input features from to evaluate as input arguments, where the dataset may be a sample of the actual training dataset.

The helper function **test_split()** is used to split the dataset by a candidate split point and **gini_index()** is used to evaluate the cost of a given split by the groups of rows created.

We can see that a list of features is created by randomly selecting feature indices and adding them to a list (called **features**), this list of features is then enumerated and specific values in the training dataset evaluated as split points.

```

1 # Select the best split point for a dataset
2 def get_split(dataset, n_features):
3     class_values = list(set(row[-1] for row in dataset))
4     b_index, b_value, b_score, b_groups = 999, 999, 999, None
5     features = list()
6     while len(features) < n_features:
7         index = randrange(len(dataset[0])-1)
8         if index not in features:
9             features.append(index)
10    for index in features:
11        for row in dataset:
12            groups = test_split(index, row[index], dataset)
13            gini = gini_index(groups, class_values)
14            if gini < b_score:
15                b_index, b_value, b_score, b_groups = index, row[index], gini, groups
16    return {'index':b_index, 'value':b_value, 'groups':b_groups}

```

```

1  # Random Forest Algorithm on Sonar Dataset
2  from random import seed
3  from random import randrange
4  from csv import reader
5  from math import sqrt
6
7  # Load a CSV file
8  def load_csv(filename):
9      dataset = list()
10     with open(filename, 'r') as file:
11         csv_reader = reader(file)
12         for row in csv_reader:
13             if not row:
14                 continue
15             dataset.append(row)
16     return dataset
17
18 # Convert string column to float
19 def str_column_to_float(dataset, column):
20     for row in dataset:
21         row[column] = float(row[column].strip())
22
23 # Convert string column to integer
24 def str_column_to_int(dataset, column):
25     class_values = [row[column] for row in dataset]
26     unique = set(class_values)
27     lookup = dict()
28     for i, value in enumerate(unique):
29         lookup[value] = i
30     for row in dataset:
31         row[column] = lookup[row[column]]
32     return lookup
33
34 # Split a dataset into k folds
35 def cross_validation_split(dataset, n_folds):
36     dataset_split = list()
37     dataset_copy = list(dataset)
38     fold_size = int(len(dataset) / n_folds)
39     for i in range(n_folds):
40         fold = list()
41         while len(fold) < fold_size:
42             index = randrange(len(dataset_copy))
43             fold.append(dataset_copy.pop(index))
44         dataset_split.append(fold)
45     return dataset_split
46
47 # Calculate accuracy percentage
48 def accuracy_metric(actual, predicted):
49     correct = 0
50     for i in range(len(actual)):
51         if actual[i] == predicted[i]:
52             correct += 1
53     return correct / float(len(actual)) * 100.0
54
55 # Evaluate an algorithm using a cross validation split
56 def evaluate_algorithm(dataset, algorithm, n_folds, *args):
57     folds = cross_validation_split(dataset, n_folds)
58     scores = list()
59     for fold in folds:
60         train_set = list(folds)
61         train_set.remove(fold)
62         train_set = sum(train_set, [])
63         test_set = list(fold)

```

```

55 # Evaluate an algorithm using a cross validation split
56 def evaluate_algorithm(dataset, algorithm, n_folds, *args):
57     folds = cross_validation_split(dataset, n_folds)
58     scores = list()
59     for fold in folds:
60         train_set = list(folds)
61         train_set.remove(fold)
62         train_set = sum(train_set, [])
63         test_set = list()
64         for row in fold:
65             row_copy = list(row)
66             test_set.append(row_copy)
67             row_copy[-1] = None
68         predicted = algorithm(train_set, test_set, *args)
69         actual = [row[-1] for row in fold]
70         accuracy = accuracy_metric(actual, predicted)
71         scores.append(accuracy)
72     return scores
73
74 # Split a dataset based on an attribute and an attribute value
75 def test_split(index, value, dataset):
76     left, right = list(), list()
77     for row in dataset:
78         if row[index] < value:
79             left.append(row)
80         else:
81             right.append(row)
82     return left, right
83
84 # Calculate the Gini index for a split dataset
85 def gini_index(groups, classes):
86     # count all samples at split point
87     n_instances = float(sum([len(group) for group in groups]))
88     # sum weighted Gini index for each group
89     gini = 0.0
90     for group in groups:
91         size = float(len(group))
92         # avoid divide by zero
93         if size == 0:
94             continue
95         score = 0.0
96         # score the group based on the score for each class
97         for class_val in classes:
98             p = [row[-1] for row in group].count(class_val) / size
99             score += p * p
100         # weight the group score by its relative size
101         gini += (1.0 - score) * (size / n_instances)
102     return gini
103
104 # Select the best split point for a dataset
105 def get_split(dataset, n_features):
106     class_values = list(set(row[-1] for row in dataset))
107     b_index, b_value, b_score, b_groups = 999, 999, 999, None
108     features = list()
109     while len(features) < n_features:
110         index = randrange(len(dataset[0])-1)
111         if index not in features:
112             features.append(index)
113     for index in features:
114         for row in dataset:
115             groups = test_split(index, row[index], dataset)
116             gini = gini_index(groups, class_values)

```

```

121 # Create a terminal node value
122 def to_terminal(group):
123     outcomes = [row[-1] for row in group]
124     return max(set(outcomes), key=outcomes.count)
125
126 # Create child splits for a node or make terminal
127 def split(node, max_depth, min_size, n_features, depth):
128     left, right = node['groups']
129     del(node['groups'])
130     # check for a no split
131     if not left or not right:
132         node['left'] = node['right'] = to_terminal(left + right)
133         return
134     # check for max depth
135     if depth >= max_depth:
136         node['left'], node['right'] = to_terminal(left), to_terminal(right)
137         return
138     # process left child
139     if len(left) <= min_size:
140         node['left'] = to_terminal(left)
141     else:
142         node['left'] = get_split(left, n_features)
143         split(node['left'], max_depth, min_size, n_features, depth+1)
144     # process right child
145     if len(right) <= min_size:
146         node['right'] = to_terminal(right)
147     else:
148         node['right'] = get_split(right, n_features)
149         split(node['right'], max_depth, min_size, n_features, depth+1)
150
151 # Build a decision tree
152 def build_tree(train, max_depth, min_size, n_features):
153     root = get_split(train, n_features)
154     split(root, max_depth, min_size, n_features, 1)
155     return root
156
157 # Make a prediction with a decision tree
158 def predict(node, row):
159     if row[node['index']] < node['value']:
160         if isinstance(node['left'], dict):
161             return predict(node['left'], row)
162         else:
163             return node['left']
164     else:
165         if isinstance(node['right'], dict):
166             return predict(node['right'], row)
167         else:
168             return node['right']
169
170 # Create a random subsample from the dataset with replacement
171 def subsample(dataset, ratio):
172     sample = list()
173     n_sample = round(len(dataset) * ratio)
174     while len(sample) < n_sample:
175         index = randrange(len(dataset))
176         sample.append(dataset[index])
177     return sample
178
179 # Make a prediction with a list of bagged trees
180 def bagging_predict(trees, row):
181     predictions = [predict(tree, row) for tree in trees]
182     return max(set(predictions), key=predictions.count)
183

```

```

155     return root
156
157 # Make a prediction with a decision tree
158 def predict(node, row):
159     if row[node['index']] < node['value']:
160         if isinstance(node['left'], dict):
161             return predict(node['left'], row)
162         else:
163             return node['left']
164     else:
165         if isinstance(node['right'], dict):
166             return predict(node['right'], row)
167         else:
168             return node['right']
169
170 # Create a random subsample from the dataset with replacement
171 def subsample(dataset, ratio):
172     sample = list()
173     n_sample = round(len(dataset) * ratio)
174     while len(sample) < n_sample:
175         index = randrange(len(dataset))
176         sample.append(dataset[index])
177     return sample
178
179 # Make a prediction with a list of bagged trees
180 def bagging_predict(trees, row):
181     predictions = [predict(tree, row) for tree in trees]
182     return max(set(predictions), key=predictions.count)
183
184 # Random Forest Algorithm
185 def random_forest(train, test, max_depth, min_size, sample_size, n_trees, n_features):
186     trees = list()
187     for i in range(n_trees):
188         sample = subsample(train, sample_size)
189         tree = build_tree(sample, max_depth, min_size, n_features)
190         trees.append(tree)
191     predictions = [bagging_predict(trees, row) for row in test]
192     return(predictions)
193
194 # Test the random forest algorithm
195 seed(2)
196 # load and prepare data
197 filename = 'sonar.all-data.csv'
198 dataset = load_csv(filename)
199 # convert string attributes to integers
200 for i in range(0, len(dataset[0])-1):
201     str_column_to_float(dataset, i)
202 # convert class column to integers
203 str_column_to_int(dataset, len(dataset[0])-1)
204 # evaluate algorithm
205 n_folds = 5
206 max_depth = 10
207 min_size = 1
208 sample_size = 1.0
209 n_features = int(sqrt(len(dataset[0])-1))
210 for n_trees in [1, 5, 10]:
211     scores = evaluate_algorithm(dataset, random_forest, n_folds, max_depth, min_size, sample_size, n_trees, n
212     print('Trees: %d' % n_trees)
213     print('Scores: %s' % scores)
214     print('Mean Accuracy: %.3f%%' % (sum(scores)/float(len(scores))))

```

URL - <https://www.kaggle.com/code/ninjaac/random-forest-from-scratch>

Random forest Regression from scratch

```
In [8]: class RandomForestRegression(RandomForest):
        """Random forest for classification task."""
        def __init__(self, max_feature, max_depth, n_trees=100, min_sample_split=2, min_impurity=1e-7):
            """
            :param max_depth: Int - Max depth of each tree.
            :param n_trees: Int - Number of trees/estimators.
            :param min_sample_split: Int - minimum samples for a node to have before going for split.
            :param min_impurity: Int - Min impurity a node can have.
            """
            self.prediction_aggrigation_calculation = self._mean_calculation

            # Initializing the trees.
            self.trees = []
            for _ in range(n_trees):
                self.trees.append(DecisionTreeRegression(min_sample_split=min_sample_split,
                                                            min_impurity=min_impurity, max_depth=max_depth))

            super().__init__(trees=self.trees, n_trees=n_trees, max_feature=max_feature,
                             prediction_aggrigation_calculation=self.prediction_aggrigation_calculation)

        def _mean_calculation(self, y_preds):
            """
```

```

"""
Find mean prediction of all tree prediction for each sample.

:param y_preds: Prediction value from number of estimators trees.
"""

# create a empty array to store the prediction.
y_pred = np.empty((y_preds.shape[0], 1))
# iterate over all the data samples.
for i, sample_predictions in enumerate(y_preds):
    y_pred[i] = np.mean(sample_predictions)

return y_pred

```

In [9]:

```

# Define the training data.
X, y = make_regression(n_samples=1000, n_features=8)

# Change the shape of the target to 1 dimensional array.
y = y[:, np.newaxis]

print("="*100)
print("Number of training data samples-----> {}".format(X.shape[0]))
print("Number of training features -----> {}".format(X.shape[1]))
print("Shape of the target value -----> {}".format(y.shape))

```

```

=====
=====
Number of training data samples-----> 1000
Number of training features -----> 8
Shape of the target value -----> (1000, 1)

```

```
# display the data.
data_y = pd.DataFrame(y)
data_y.head()
```

Out[10]:

	0
0	-261.609902
1	-66.627855
2	50.562104
3	5.719164
4	-33.019911

In [11]:

```
#define the parameters
print("="*100)
random_forest_reg = RandomForestRegression(n_trees=3, max_feature=3, min_sample_split=2, max_dept
h=45)

# Train the model.
random_forest_reg.train(X, y)

# Predict the values.
y_pred = random_forest_reg.predict(X)
#Root mean square error.
score = r2_score(y, y_pred)
print("The r2_score of the trained model", score)
```

```
=====
=====
```

1	0
2	1
3	1
4	1

In [7]:

```
#define the parameters

param = {
    "n_neighbours" : 5
}

print("=*100)
random_forest_cla = RandomForestClassifier(n_trees=3, max_feature=2, min_sample_split=2, max_dept
h=45)

# Train the model.
random_forest_cla.train(X, y)

# Predict the values.
y_pred = random_forest_cla.predict(X)

#calculate accuracy.
acc = np.sum(y==y_pred)/X.shape[0]
print("=*100)
print("Accuracy of the prediction is {}".format(acc))
```

URL - <https://www.analyticsvidhya.com/blog/2018/12/building-a-random-forest-from-scratch-understanding-real-world-data-products-ml-for-programmers-part-3/>

Note: Steps 1 to 6 have been covered in the previous article.

Step 1: Importing the basic libraries.

```
%load_ext autoreload
%autoreload 2
%matplotlib inline

from fastai.imports import *
from fastai.structured import *
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier

from IPython.display import display
from sklearn import metrics
```

Step 2: Read the data and split into train and validation sets.

```
PATH = "data/bulldozers/"

df_raw = pd.read_feather('tmp/bulldozers-raw')
df_trn, y_trn, nas = proc_df(df_raw, 'SalePrice')

def split_vals(a,n): return a[:n], a[n:]
n_valid = 12000
n_trn = len(df_trn)-n_valid

X_train, X_valid = split_vals(df_trn, n_trn)
y_train, y_valid = split_vals(y_trn, n_trn)

raw_train, raw_valid = split_vals(df_raw, n_trn)
```

Step 3: Take a subset of data to start with.

As I previously mentioned, we'll take smaller steps, so here we pick only two features and work with them. If that works well, we can complete the model by taking all the features.

```
x_sub = X_train[['YearMade', 'MachineHoursCurrentMeter']]
```

Step 4: Define the set of inputs:

- A set of features – x
- A target variable – y
- Number of trees in the random forest – n_trees
- A variable to define the sample size – sample_sz
- A variable for minimum leaf size – min_leaf
- A random seed for testing

```
def __init__(self, x, y, n_trees, sample_sz, min_leaf=5):
    np.random.seed(42)
    self.x, self.y, self.sample_sz, self.min_leaf = x, y, sample_sz, min_leaf
    self.trees = [self.create_tree() for i in range(n_trees)]
```

Step 5: Define a function that uses a sample of data (with replacement) and creates a decision tree over the same.

```
def create_tree(self):
    rnd_idx = np.random.permutation(len(self.y))[:self.sample_sz]
    return DecisionTree(self.x.iloc[rnd_idx], self.y[rnd_idx], min_leaf=self.min_leaf)
```

Step 6: Create a predict function. The mean of the predicted value from each tree for a particular row is returned as the final prediction.

```
def predict(self, x):
    return np.mean([t.predict(x) for t in self.trees], axis=0)
```

Combining all the above functions, we can create a class *TreeEnsemble*.

```
class TreeEnsemble():
    def __init__(self, x, y, n_trees, sample_sz, min_leaf=5):
        np.random.seed(42)
        self.x, self.y, self.sample_sz, self.min_leaf = x, y, sample_sz, min_leaf
        self.trees = [self.create_tree() for i in range(n_trees)]

    def create_tree(self):
        rnd_idx = np.random.permutation(len(self.y))[:self.sample_sz]
        return DecisionTree(self.x.iloc[rnd_idx], self.y[rnd_idx], min_leaf=self.min_leaf)

    def predict(self, x):
        return np.mean([t.predict(x) for t in self.trees], axis=0)
```

Step 7: Create a class *DecisionTree*. We call *DecisionTree* in the function *create_tree*, so let's define it here. A decision tree would have a set of independent variables, a target variable, and the index values. For now, we have created only one decision tree (we can make it recursive later).

```
class DecisionTree():
    def __init__(self, x, y, idxs=None, min_leaf=5):
        if idxs is None: idxs=np.arange(len(y))
        self.x,self.y,self.idxs,self.min_leaf = x,y,idxs,min_leaf ##define x,y,index and minimum leaf size
        self.n,self.c = len(idxs), x.shape[1] ##number of rows and columns
        self.val = np.mean(y[idxs])
        self.score = float('inf')
        self.find_varsplit()
```

Step 8: Determine the best split point. For every column, we use the function *find_better_split* to identify a split point and then return the column name, value, and score for the split.

```
    def find_varsplit(self):
        for i in range(self.c): self.find_better_split(i) #check for each column in the dataset

def find_varsplit(self):
    for i in range(self.c): self.find_better_split(i)

@property
def split_name(self): return self.x.columns[self.var_idx]

@property
def split_col(self):
    return self.x.values[self.idxs,self.var_idx]

@property
def is_leaf(self): return self.score == float('inf')

def __repr__(self):
    s = f'n: {self.n}; val:{self.val}'
```

```

@property
def is_leaf(self): return self.score == float('inf')

def __repr__(self):
    s = f'n: {self.n}; val:{self.val}'
    if not self.is_leaf:
        s += f'; score:{self.score}; split:{self.split}; var:
            {self.split_name}'
    return s

```

Step 9: Build our first model with 10 trees, a sample size of 1,000 and minimum leaf as 3.

```

m = TreeEnsemble(X_train, y_train, n_trees=10, sample_sz=1000,min_leaf=3)
m.trees[0]

```

For the first tree, the results are:

```
n: 1000; val:10.079014121552744
```

Let's now fill the block that we left above in Step 8 – *find_better_split*. This is by far the most complicated part of the code to understand, but Jeremy has explained using a simple example in excel. I will explain it in an intuitive manner here.

For each variable, we split the points to the left and right node, and check the score for every value. The idea is to find a split point where we are able to separate more similar points together.

Consider the following example: we have two columns – an independent variable which we try to split on, and a binary target variable.

We can take a weighted average of this value. Similarly, we can calculate for a split at 4, 6, 1 and so on. Let's put this into code:

```
def find_better_split(self, var_idx):
    x,y = self.x.values[self.idxs,var_idx], self.y[self.idxs]

    for i in range(self.n):
        lhs = x<=x[i]
        rhs = x>x[i]
        if rhs.sum()<self.min_leaf or lhs.sum()<self.min_leaf: continue
        lhs_std = y[lhs].std()
        rhs_std = y[rhs].std()
        curr_score = lhs_std*lhs.sum() + rhs_std*rhs.sum()
        if curr_score<self.score:
            self.var_idx,self.score,self.split = var_idx,curr_score,x[i]
```

If we try to print the results of the function for both the columns individually, we get the below result:

```
find_better_split(tree,1)
tree

n: 1000; val:10.079014121552744; score:681.0184057251435; split:3744.0; var:MachineHoursCurrentMeter

find_better_split(tree,0);
tree

n: 1000; val:10.079014121552744; score:658.5510186055949; split:1974.0; var:YearMade
```


Looks like *YearMade* 1974 is a better split point.

Step 10: Compare with the *scikit-learn* random forest. But keep in mind that there's a tricky aspect here. While comparing the two models, both of them should have the same input. So let's store the input that we have used in the random forest we just built.

```
ens = TreeEnsemble(x_sub, y_train, 1, 1000)
tree = ens.trees[0]
x_samp, y_samp = tree.x, tree.y
```

And now we build a model on this subset:

```
m = RandomForestRegressor(n_estimators=1, max_depth=1, bootstrap=False)
m.fit(x_samp, y_samp)
draw_tree(m.estimators_[0], x_samp, precision=2)
```



We see that the split here is on the column *YearMade* at year 1974.5, very similar to the results of our model. Not bad!

Step 11: There's a problem with the code we have seen so far – can you recognize what it is? We need to optimise it! In the current format, we are checking the score for the split at each row, aka we are checking for a single value multiple times. Have a look at the example we used earlier:

Feature	Target
3	0
4	1
1	0
6	1
4	0
1	0

The function will check the split points 4 and 1 twice because it actually works row-wise. It's a good idea to optimise our code so as to reduce the computation time (not everyone has a top machine!). The idea is to sort column-wise and then check the score after the split for unique values only.

```
sort_idx = np.argsort(x)
sort_y, sort_x = y[sort_idx], x[sort_idx]
```

Also, to calculate the standard deviation, we define the following function:

```
def std_agg(cnt, s1, s2): return math.sqrt((s2/cnt) - (s1/cnt)**2)
```

We will need to keep a track of the count of data points on each side of the split along with the sum of square of the values. So we initialize the variables *rhs_cnt*, *lhs_cnt*, *rjs_sum2* and *lhs_sum2*.

```
rhs_cnt, rhs_sum, rhs_sum2 = self.n, sort_y.sum(), (sort_y**2).sum()
lhs_cnt, lhs_sum, lhs_sum2 = 0, 0, 0
```

Adding all this up, the code looks like this:

```
tree = TreeEnsemble(x_sub, y_train, 1, 1000).trees[0]
```

Adding all this up, the code looks like this:

```
tree = TreeEnsemble(x_sub, y_train, 1, 1000).trees[0]
def std_agg(cnt, s1, s2): return math.sqrt((s2/cnt) - (s1/cnt)**2)
def find_better_split_foo(self, var_idx):
    x,y = self.x.values[self.idxs,var_idx], self.y[self.idxs]

    sort_idx = np.argsort(x)
    sort_y,sort_x = y[sort_idx], x[sort_idx]

    rhs_cnt,rhs_sum,rhs_sum2 = self.n, sort_y.sum(), (sort_y**2).sum()
    lhs_cnt,lhs_sum,lhs_sum2 = 0,0.,0.

    for i in range(0,self.n-self.min_leaf-1):
        xi,yi = sort_x[i],sort_y[i]
        lhs_cnt += 1; rhs_cnt -= 1
        lhs_sum += yi; rhs_sum -= yi
        lhs_sum2 += yi**2; rhs_sum2 -= yi**2
        if i<self.min_leaf or xi==sort_x[i+1]:
            continue

        lhs_std = std_agg(lhs_cnt, lhs_sum, lhs_sum2)
        rhs_std = std_agg(rhs_cnt, rhs_sum, rhs_sum2)
        curr_score = lhs_std*lhs_cnt + rhs_std*rhs_cnt
        if curr_score<self.score:
            self.var_idx,self.score,self.split = var_idx,curr_score,xi
```

Ideally, this function should give the same results. Let's check:

```
%timeit find_better_split_foo(tree,1)
tree

2.2 ms ± 148 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
n: 1000; val:10.079014121552744; score:658.5510186055565; split:1974.0; var:YearMade
```

```
DecisionTree.find_better_split = find_better_split_too
```

Step 12: Build our tree with more than one split.

In step 10, we compared the first level of our model with the scikit-learn random forest model. We will now create a full tree (which splits on both our features) and compare it again. Right now our *find_varsplit* function looks like this:

```
def find_varsplit(self):
    for i in range(self.c): self.find_better_split(i)
```

where we have defined *find_better_split* separately. We will update this function, so that it automatically checks for the leaf node and stores a list of indices in lhs and rhs after the split.

```
def find_varsplit(self):
    for i in range(self.c): self.find_better_split(i)
    if self.is_leaf: return
    x = self.split_col
    lhs = np.nonzero(x<=self.split)[0]
    rhs = np.nonzero(x>self.split)[0]
    self.lhs = DecisionTree(self.x, self.y, self.idxs[lhs])
    self.rhs = DecisionTree(self.x, self.y, self.idxs[rhs])
```

```
DecisionTree.find_varsplit = find_varsplit
```

We will again compare both the models. Previously, the *max_depth* was restricted to 1, and we will make it 2 here (we have only two features for now):

```
m = RandomForestRegressor(n_estimators=1, max_depth=2, bootstrap=False)
m.fit(x_samp, y_samp)
draw_tree(m.estimators_[0], x_samp, precision=2)
```

Step 13: Calculate the final predictions. We have called a predict function in *TreeEnsemble* that returns the prediction for each row:

```
def predict(self, x):
    return np.array([self.predict_row(xi) for xi in x])

def predict_row(self, xi):
    if self.is_leaf: return self.val
    t = self.lhs if xi[self.var_idx]<=self.split else self.rhs
    return t.predict_row(xi)
DecisionTree.predict_row = predict_row
```

With this, we have completed building our own random forest model. Let's plot the predictions on the validation set:

```
cols = ['MachineID', 'YearMade', 'MachineHoursCurrentMeter',
        'ProductSize', 'Enclosure', 'Coupler_System', 'saleYear']

%time tree = TreeEnsemble(X_train[cols], y_train, 1, 1000).trees[0]
x_samp, y_samp = tree.x, tree.y

CPU times: user 288 ms, sys: 12 ms, total: 300 ms
Wall time: 297 ms

plt.scatter(preds, y_valid, alpha=0.05)
```

URL - <https://www.pycodemates.com/2022/10/implementing-SVM-from-scratch-in-python.html>

Implementing SVM from scratch in python

Writing the SVM class

```
# svm.py

import numpy as np

class SVM:

    def __init__(self, C = 1.0):
        # C = error term
        self.C = C
        self.w = 0
        self.b = 0
```

Hinge Loss calculation

```
# Hinge Loss Function / Calculation
def hingeloss(self, w, b, x, y):
    # Regularizer term
    reg = 0.5 * (w * w)

    for i in range(x.shape[0]):
        # Optimization term
        opt_term = y[i] * ((np.dot(w, x[i])) + b)

        # calculating loss
        loss = reg + self.C * max(0, 1-opt_term)
    return loss[0][0]
```

Let's understand what's happening here. First, we are calculating the value of the **regularizer** term and assign it to variable **reg**. Then iterating over the number of samples and calculating the **loss** using the **optimization** function.

Gradient Descent optimization

Now let's create the gradient descent function inside the fit method in order to get the best parameters **w** and **b**.

```
def fit(self, X, Y, batch_size=100, learning_rate=0.001, epochs=1000):  
    # The number of features in X  
    number_of_features = X.shape[1]  
  
    # The number of Samples in X  
    number_of_samples = X.shape[0]  
  
    c = self.C  
  
    # Creating ids from 0 to number_of_samples - 1  
    ids = np.arange(number_of_samples)  
  
    # Shuffling the samples randomly  
    np.random.shuffle(ids)  
  
    # creating an array of zeros  
    w = np.zeros((1, number_of_features))  
    b = 0  
    losses = []  
  
    # Gradient Descent Logic  
    for i in range(epochs):  
        # Calculating the Hinge Loss  
        l = self.hingeloss(w, b, X, Y)  
  
        # Appending all losses  
        losses.append(l)  
  
        # Starting from 0 to the number of samples with batch_size as interval  
        for batch_initial in range(0, number_of_samples, batch_size):  
            gradw = 0
```

```

# Calculating the Hinge Loss
for i in range(epochs):
    # Calculating the Hinge Loss
    l = self.hingeloss(w, b, X, Y)

    # Appending all losses
    losses.append(l)

    # Starting from 0 to the number of samples with batch_size as interval
    for batch_initial in range(0, number_of_samples, batch_size):
        gradw = 0
        gradb = 0

        for j in range(batch_initial, batch_initial + batch_size):
            if j < number_of_samples:
                x = ids[j]
                ti = Y[x] * (np.dot(w, X[x].T) + b)

                if ti > 1:
                    gradw += 0
                    gradb += 0
                else:
                    # Calculating the gradients

                    #w.r.t w
                    gradw += c * Y[x] * X[x]
                    # w.r.t b
                    gradb += c * Y[x]

            # Updating weights and bias
            w = w - learning_rate * w + learning_rate * gradw
            b = b + learning_rate * gradb

self.w = w
self.b = b

return self.w, self.b, losses

```

Splitting the dataset

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5, random_state=42)
```

Training the SVM model

```
svm = SVM()

w, b, losses = svm.fit(X_train, y_train)
```

Making predictions and testing accuracy

```
prediction = svm.predict(X_test)

# Loss value
lss = losses.pop()

print("Loss:", lss)
print("Prediction:", prediction)
print("Accuracy:", accuracy_score(prediction, y_test))
print("w, b:", [w, b])

----

Loss: 0.0991126738798482
Prediction: [-1.  1. -1. -1.  1.  1.  1.  1. -1.  1. -1. -1.  1.  1. -1. -1.  1. -1.
  1. -1.  1.  1. -1.  1.  1.  1. -1. -1. -1. -1.  1. -1. -1.  1.  1. -1.
  1. -1. -1.  1.  1.  1.  1.  1.  1. -1.  1.  1.  1.  1.]
Accuracy: 1.0
w, b: [array([[0.44477983, 0.15109913]]), 0.057000000000000004]
```

URL – https://www.python-engineer.com/courses/mlfromscratch/07_svm/

Implementation

```
import numpy as np

class SVM:

    def __init__(self, learning_rate=0.001, lambda_param=0.01, n_iters=1000):
        self.lr = learning_rate
        self.lambda_param = lambda_param
        self.n_iters = n_iters
        self.w = None
        self.b = None

    def fit(self, X, y):
        n_samples, n_features = X.shape

        y_ = np.where(y <= 0, -1, 1)

        self.w = np.zeros(n_features)
        self.b = 0

        for _ in range(self.n_iters):
            for idx, x_i in enumerate(X):
                condition = y_[idx] * (np.dot(x_i, self.w) - self.b) >= 1
                if condition:
                    self.w -= self.lr * (2 * self.lambda_param * self.w)
                else:
                    self.w -= self.lr * (2 * self.lambda_param * self.w - np.dot(x_i, y_[idx]))
                    self.b -= self.lr * y_[idx]

    def predict(self, X):
        approx = np.dot(X, self.w) - self.b
        return np.sign(approx)
```

URL – <https://towardsdatascience.com/implement-multiclass-svm-from-scratch-in-python-b141e43dc084>

Defining Kernels and SVM Hyperparameters

We start by defining the three kernels using their respective functions

$K(x, x')$	Linear	Polynomial	Gaussian (Radial Basis Fun.)
	$x^T x'$	$\phi(x) = (1 + x^T x')^Q$	$e^{-\gamma \ x - x'\ ^2}$

Degree of polynomial (Q) and RBF γ are hyperparameters (decided by the user)

```
class SVM:
    linear = lambda x, x', c=0: x @ x'.T
    polynomial = lambda x, x', Q=5: (1 + x @ x'.T)**Q
    rbf = lambda x, x', gamma=10: np.exp(-gamma*distance.cdist(x, x', 'sqeuclidean'))
    kernel_funs = {'linear': linear, 'polynomial': polynomial, 'rbf': rbf}
```

For consistency with other kernels, the linear kernel takes an extra useless hyperparameter. As obvious, `kernel_funs` takes a string for the kernel and returns the corresponding kernel function.

Now let's carry on by defining the constructor:

```
class SVM:
    linear = lambda x, x', c=0: x @ x'.T
    polynomial = lambda x, x', Q=5: (1 + x @ x'.T)**Q
    rbf = lambda x, x', gamma=10: np.exp(-gamma*distance.cdist(x, x', 'sqeuclidean'))
    kernel_funs = {'linear': linear, 'polynomial': polynomial, 'rbf': rbf}

    def __init__(self, kernel='rbf', C=1, k=2):
        # set the hyperparameters
        self.kernel = kernel
        self.C = C
        self.k = k
```

Now let's carry on by defining the constructor:

```
class SVM:
    linear = lambda x, x', c=0: x @ x'.T
    polynomial = lambda x, x', Q=5: (1 + x @ x'.T)**Q
    rbf = lambda x, x', gamma=10: np.exp(-gamma*distance.cdist(x, x', 'sqeuclidean'))
    kernel_funs = {'linear': linear, 'polynomial': polynomial, 'rbf': rbf}

    def __init__(self, kernel='rbf', C=1, k=2):
        # set the hyperparameters
        self.kernel_str = kernel
        self.kernel = SVM.kernel_funs[kernel]
        self.C = C                # regularization parameter
        self.k = k                # kernel parameter

        # training data and support vectors (set later)
        self.X, y = None, None
        self.as = None

        # for multi-class classification (set later)
        self.multiclass = False
        self.clfs = []
```

The SVM has three main hyperparameters, the kernel (we store the string given and the corresponding kernel function), the regularization parameter C and the kernel hyperparameter (to be passed to the kernel function); it represents Q for the polynomial kernel and γ for the RBF kernel.

Define the Fit Method

To extend this class with `fit` and `predict` functions in separate cells, we will define the following function and use it as a decorator later:

By this, we can write the following fit function:

```
@SVMClass
def fit(self, X, y, eval_train=False):
    # if more than two unique labels, call the multiclass version
    if len(np.unique(y)) > 2:
        self.multiclass = True
        return self.multi_fit(X, y, eval_train)

    # if labels given in {0,1} change it to {-1,1}
    if set(np.unique(y)) == {0, 1}: y[y == 0] = -1

    # ensure y is a Nx1 column vector (needed by CVXOPT)
    self.y = y.reshape(-1, 1).astype(np.double) # Has to be a column vector
    self.X = X
    N = X.shape[0] # Number of points

    # compute the kernel over all possible pairs of (x, x') in the data
    # by Numpy's vectorization this yields the matrix K
    self.K = self.kernel(X, X, self.k)

    ### Set up optimization parameters
    # For  $\frac{1}{2} x^T P x + q^T x$ 
    P = cvxopt.matrix(self.y @ self.y.T * self.K)
    q = cvxopt.matrix(-np.ones((N, 1)))

    # For  $Ax = b$ 
    A = cvxopt.matrix(self.y.T)
    b = cvxopt.matrix(np.zeros(1))

    # For  $Gx \leq h$ 
    G = cvxopt.matrix(np.vstack((-np.identity(N),
                                  np.identity(N))))
    h = cvxopt.matrix(np.vstack((np.zeros((N,1)),
                                  np.ones((N,1)) * self.C)))

    # Solve
    cvxopt.solvers.options['show_progress'] = False
    sol = cvxopt.solvers.qp(P, q, G, h, A, b)
    self.alpha = np.array(sol["x"]) # our solution
```

```

np.ones((N,1)) * self.C)))

# Solve
cvxopt.solvers.options['show_progress'] = False
sol = cvxopt.solvers.qp(P, q, G, h, A, b)
self.αs = np.array(sol["x"])          # our solution

# a Boolean array that flags points which are support vectors
self.is_sv = ((self.αs-1e-3 > 0)&(self.αs <= self.C)).squeeze()
# an index of some margin support vector
self.margin_sv = np.argmax((0 < self.αs-1e-3)&(self.αs < self.C-1e-3))

if eval_train:
    print(f"Finished training with accuracy{self.evaluate(X, y)}")

```

```

@SVMClass
def predict(self, X_t):
    if self.multiclass: return self.multi_predict(X_t)
    # compute (x, y)
    x, y, = self.X[self.margin_sv, np.newaxis], self.y[self.margin_sv]
    # find support vectors
    αs, y, X= self.αs[self.is_sv], self.y[self.is_sv], self.X[self.is_sv]
    # compute the second term
    b = y, - np.sum(αs * y * self.kernel(X, x, self.k), axis=0)
    # compute the score
    score = np.sum(αs * y * self.kernel(X, X_t, self.k), axis=0) + b
    return np.sign(score).astype(int), score

```

That's it. We can also implement an `evaluate` method to compute the accuracy (used in fit above).

```

@SVMClass
def evaluate(self, X,y):
    outputs, _ = self.predict(X)
    accuracy = np.sum(outputs == y) / len(y)
    return round(accuracy, 2)

```

Test the Implementation

```

from sklearn.datasets import make_classification
import numpy as np

# Load the dataset
np.random.seed(1)
X, y = make_classification(n_samples=2500, n_features=5,
                          n_redundant=0, n_informative=5,
                          n_classes=2, class_sep=0.3)

# Test Implemented SVM
svm = SVM(kernel='rbf', k=1)
svm.fit(X, y, eval_train=True)

y_pred, _ = svm.predict(X)
print(f"Accuracy: {np.sum(y==y_pred)/y.shape[0]}") #0.9108

# Test with Scikit
from sklearn.svm import SVC
clf = SVC(kernel='rbf', C=1, gamma=1)
clf.fit(X, y)
y_pred = clf.predict(X)
print(f"Accuracy: {sum(y==y_pred)/y.shape[0]}") #0.9108

```

Generalizing Fit to Multiclass

```

@SVMClass
def multi_fit(self, X, y, eval_train=False):
    self.k = len(np.unique(y)) # number of classes
    # for each pair of classes
    for i in range(self.k):
        # get the data for the pair
        Xs, Ys = X, copy.copy(y)
        # change the labels to -1 and 1
        Ys[Ys!=i], Ys[Ys==i] = -1, +1
        # fit the classifier
        clf = SVM(kernel=self.kernel_str, C=self.C, k=self.k)
        clf.fit(Xs, Ys)
        # save the classifier
        self.clfs.append(clf)
    if eval_train:
        print(f"Finished training with accuracy {self.evaluate(X, y)}")

```

Generalizing Predict to Multiclass

Then to perform prediction on a new example, we choose the class for which the corresponding classifier is most confident (has the highest score)

```
@SVMClass
def multi_predict(self, X):
    # get the predictions from all classifiers
    N = X.shape[0]
    preds = np.zeros((N, self.k))
    for i, clf in enumerate(self.clfes):
        _, preds[:, i] = clf.predict(X)

    # get the argmax and the corresponding score
    return np.argmax(preds, axis=1), np.max(preds, axis=1)
```

Testing the Implementation

```
from sklearn.datasets import make_classification
import numpy as np

# Load the dataset
np.random.seed(1)
X, y = make_classification(n_samples=500, n_features=2,
                           n_redundant=0, n_informative=2,
                           n_classes=4, n_clusters_per_class=1,
                           class_sep=0.3)

# Test SVM
svm = SVM(kernel='rbf', k=4)
svm.fit(X, y, eval_train=True)

y_pred = svm.predict(X)
print(f"Accuracy: {np.sum(y==y_pred)/y.shape[0]}") # 0.65

# Test with Scikit
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC

clf = OneVsRestClassifier(SVC(kernel='rbf', C=1, gamma=4)).fit(X, y)
y_pred = clf.predict(X)
print(f"Accuracy: {sum(y==y_pred)/y.shape[0]}") # 0.65
```

URL - <https://python.plainenglish.io/support-vector-machines-svm-from-scratch-to-libraries-990086a666d4>

```
import numpy as np
class SVM:
    def __init__(self, learning_rate=0.001, lambda_param=0.01, n_iters=1000):
        self.lr = learning_rate
        self.lambda_param = lambda_param
        self.n_iters = n_iters
        self.w = None
        self.b = None
    def fit(self, X, y):
        n_samples, n_features = X.shape
        y_ = np.where(y <= 0, -1, 1) # Convert labels to -1 and 1
        self.w = np.zeros(n_features)
        self.b = 0
        for _ in range(self.n_iters):
            for idx, x_i in enumerate(X):
                condition = y_[idx] * (np.dot(x_i, self.w) - self.b) >= 1
                if condition:
                    self.w -= self.lr * (2 * self.lambda_param * self.w)
                else:
                    self.w -= self.lr * (2 * self.lambda_param * self.w - np.dot(x_i, y_[idx]))
                    self.b -= self.lr * y_[idx]
    def predict(self, X):
        approx = np.dot(X, self.w) - self.b
        return np.sign(approx)
# Example usage:
# svm = SVM()
# svm.fit(X_train, y_train)
# predictions = svm.predict(X_test)
```

While understanding the core principles is crucial, practical implementations often leverage established libraries like Scikit-learn in Python. These libraries provide optimized versions of SVMs and support various kernels, making them suitable for a wide range of applications.

```
# Using SVM from Scikit-learn
from sklearn import svm
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Generate some data
X, y = make_classification(n_samples=100, n_features=20, random_state=42)

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random=

# Create an SVM classifier with a linear kernel
clf = svm.SVC(kernel='linear')

clf.fit(X_train, y_train)

# Predict on the test set and calculate accuracy
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")

## try it on ya pc ##
```

+

Create

🏠

Home

🏆

Competitions

📁

Datasets

👤

Models

📄

Code

💬

Discussions

📅

Learn

⋮

More

📁

Your Work

VIEWED

Weather Type Classifi...

Random Forest from s...

SVM from scratch

Data Science Salaries ...

Loan-Approval-Predic...

🔍

Search

Weather Type Classification

Forecast with Precision: Simulated Data for Predicting Weather Types

Data Card

Code (79)

Discussion (2)

Suggestions (0)

About Dataset

Description

Variables

Usability

10.00

License

Other (specified in description)

Expected update frequency

Annually

Tags

Beginner

Data Visualization

Classification

Exploratory Data Analysis

Weather and Climate

📄

Description

This dataset is synthetically generated to mimic weather data for classification tasks. It includes various weather-related features and categorizes the weather into four types: Rainy, Sunny, Cloudy, and Snowy. This dataset is designed for practicing classification algorithms, data preprocessing, and outlier detection methods.

📄

Variables

- Temperature (numeric):** The temperature in degrees Celsius, ranging from extreme cold to extreme heat.
- Humidity (numeric):** The humidity percentage, including values above 100% to introduce outliers.
- Wind Speed (numeric):** The wind speed in kilometers per hour, with a range including unrealistically high values.
- Precipitation (%) (numeric):** The precipitation percentage, including outlier values.
- Cloud Cover (categorical):** The cloud cover description.
- Atmospheric Pressure (numeric):** The atmospheric pressure in hPa, covering a wide range.
- UV Index (numeric):** The UV index, indicating the strength of ultraviolet radiation.
- Season (categorical):** The season during which the data was recorded.
- Visibility (km) (numeric):** The visibility in kilometers, including very low or very high values.
- Location (categorical):** The type of location where the data was recorded.
- Weather Type (categorical):** The target variable for classification, indicating the weather type.

+

Create

🏠

Home

🏆

Competitions

📁

Datasets

🤖

Models

🔗

Code

💬

Discussions

📖

Learn

⋮

More

📁

Your Work

VIEWED

📄

Loan-Approval-Predic...

📄

SVM From Scratch

📄

Weather Type Classifi...

📄

Random Forest from s...

📄

SVM from scratch

🔍

Search

👤

YALA KHORASANI · UPDATED 2 MONTHS AGO

524

New Notebook

Download

📄

Mobile Device Usage and User Behavior Dataset

Analyzing Mobile Usage Patterns and User Behavior Classification Across Devices

Data Card

Code (110)

Discussion (1)

Suggestions (0)

About Dataset

This dataset provides a comprehensive analysis of mobile device usage patterns and user behavior classification. It contains 700 samples of user data, including metrics such as app usage time, screen-on time, battery drain, and data consumption. Each entry is categorized into one of five user behavior classes, ranging from light to extreme usage, allowing for insightful analysis and modeling.

Key Features:

- User ID: Unique identifier for each user.
- Device Model: Model of the user's smartphone.
- Operating System: The OS of the device (iOS or Android).
- App Usage Time: Daily time spent on mobile applications, measured in minutes.
- Screen On Time: Average hours per day the screen is active.
- Battery Drain: Daily battery consumption in mAh.
- Number of Apps Installed: Total apps available on the device.
- Data Usage: Daily mobile data consumption in megabytes.
- Age: Age of the user.
- Gender: Gender of the user (Male or Female).

Usability

10.00

License

Apache 2.0

Expected update frequency

Annually

Tags

Computer Science

Data Visualization

Classification

NLP

Mobile and Wireless

9.3 Notebook

 jupyter Data_Mining_Final_Project Last Checkpoint: a few seconds ago (autosaved)  Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

         Code 

Data Mining Final Project

Library Imports

```
In [ ]: import pandas as pd
import numpy as np
from typing import Dict, List, Tuple, Union, Optional
```

Defining Functions and Classes

One-Hot Encoding Function

```
In [ ]: # Manual one-hot encoding for categorical features
def one_hot_encode(df, column):
    unique_values = df[column].unique()
    for value in unique_values:
        df[f"{column}_{value}"] = np.where(df[column] == value, 1, 0)
    df.drop(columns=[column], inplace=True)
    return df
```

Standardization Function

```
In [ ]: # Custom standardization function
def standardize(X):
    mean = X.mean(axis=0)
    std = X.std(axis=0)
    X_standardized = (X - mean) / std
    return X_standardized
```

Train Test Split Function

```
In [ ]: def train_test_split(X, y, test_size=0.3, random_seed=42):
    np.random.seed(random_seed)
    indices = np.arange(X.shape[0])
    np.random.shuffle(indices)
```

Train Test Split Function

```
In [ ]: def train_test_split(X, y, test_size=0.3, random_seed=42):
        np.random.seed(random_seed)
        indices = np.arange(X.shape[0])
        np.random.shuffle(indices)

        test_size = int(test_size * X.shape[0])
        test_indices = indices[:test_size]
        train_indices = indices[test_size:]

        X_train, X_test = X[train_indices], X[test_indices]
        y_train, y_test = y[train_indices], y[test_indices]

        return X_train, X_test, y_train, y_test
```

Majority Vote Function

```
In [ ]: def majority_vote(predictions):
        # Transpose the predictions so that each row represents predictions for one sample
        predictions = list(map(list, zip(*predictions)))

        majority_votes = []
        for preds in predictions:
            # Create a dictionary to count occurrences of each class
            vote_counts = {}
            for pred in preds:
                if pred in vote_counts:
                    vote_counts[pred] += 1
                else:
                    vote_counts[pred] = 1
            # Find the class with the maximum votes
            majority_class = max(vote_counts, key=vote_counts.get)
            majority_votes.append(majority_class)

        return majority_votes
```

Decision Tree Class

```
In [ ]: class DecisionTree:
        def __init__(self, max_depth=None, min_samples_split=2, min_samples_leaf=1):
            self.max_depth = max_depth
            self.min_samples_split = min_samples_split
            self.min_samples_leaf = min_samples_leaf
            self.tree = None

        def fit(self, X, y, depth=0):
            # Check stopping criteria
            if (self.max_depth is not None and depth >= self.max_depth) or \
                len(y) < self.min_samples_split or \
                len(set(y)) == 1:
                self.tree = self._majority_class(y)
                return

            # Find the best split
            best_feature, best_threshold = self._best_split(X, y)
            if best_feature is None:
                self.tree = self._majority_class(y)
                return

            # Create child nodes by splitting the dataset
            left_indices = X[:, best_feature] < best_threshold
            right_indices = ~left_indices

            # Check minimum samples leaf criterion
            if np.sum(left_indices) < self.min_samples_leaf or np.sum(right_indices) < self.min_samples_leaf:
                self.tree = self._majority_class(y)
                return

            left_child = DecisionTree(
                max_depth=self.max_depth,
                min_samples_split=self.min_samples_split,
                min_samples_leaf=self.min_samples_leaf
            )
            right_child = DecisionTree(
                max_depth=self.max_depth,
                min_samples_split=self.min_samples_split,
                min_samples_leaf=self.min_samples_leaf
            )

            # Recursive fitting
            left_child.fit(X[left_indices], y[left_indices], depth + 1)
            right_child.fit(X[right_indices], y[right_indices], depth + 1)

            # Store tree as a dictionary
```

```

# Store tree as a dictionary
self.tree = {
    'feature': best_feature,
    'threshold': best_threshold,
    'left': left_child,
    'right': right_child
}

def _best_split(self, X, y):
    best_gini = float('inf')
    best_feature = None
    best_threshold = None

    # If max_features is specified, randomly select features
    n_features = X.shape[1]

    for feature in range(n_features):
        thresholds = np.unique(X[:, feature])
        for threshold in thresholds:
            left_indices = X[:, feature] < threshold
            right_indices = ~left_indices

            if np.sum(left_indices) < self.min_samples_leaf or \
               np.sum(right_indices) < self.min_samples_leaf:
                continue

            gini = self._gini_impurity(y[left_indices], y[right_indices])
            if gini < best_gini:
                best_gini = gini
                best_feature = feature
                best_threshold = threshold

    return best_feature, best_threshold

def _gini_impurity(self, left_y, right_y):
    def gini(y):
        m = len(y)
        if m == 0:
            return 0
        return 1.0 - sum((np.sum(y == c) / m) ** 2 for c in np.unique(y))

    left_size = len(left_y)
    right_size = len(right_y)
    total_size = left_size + right_size

    if total_size == 0:
        return float('inf')

```

```

        return float('inf')

    weighted_gini = (left_size / total_size) * gini(left_y) + \
        (right_size / total_size) * gini(right_y)
    return weighted_gini

def predict(self, X):
    return np.array([self._predict_sample(sample) for sample in X])

def _predict_sample(self, sample):
    node = self.tree
    while isinstance(node, dict):
        if sample[node['feature']] < node['threshold']:
            node = node['left'].tree
        else:
            node = node['right'].tree
    return node

def _majority_class(self, y):
    return max(set(y), key=list(y).count)

```

Random Forest Class

```

In [ ]: class RandomForest:
    def __init__(self, n_estimators=100, max_depth=None, min_samples_split=2,
                  min_samples_leaf=1, max_features='sqrt'):
        self.n_estimators = n_estimators
        self.max_depth = max_depth
        self.min_samples_split = min_samples_split
        self.min_samples_leaf = min_samples_leaf
        self.max_features = max_features
        self.trees = []

    def _get_max_features(self, n_features):
        if isinstance(self.max_features, str):
            if self.max_features == 'sqrt':
                return max(1, int(np.sqrt(n_features)))
            elif self.max_features == 'log2':
                return max(1, int(np.log2(n_features)))
        elif isinstance(self.max_features, float):
            return max(1, int(self.max_features * n_features))
        elif isinstance(self.max_features, int):
            return min(self.max_features, n_features)
        return n_features

    def _bootstrap_sample(self, X, y):
        n_samples = X.shape[0]
        indices = np.random.choice(n_samples, n_samples, replace=True)

```

```

def _bootstrap_sample(self, X, y):
    n_samples = X.shape[0]
    indices = np.random.choice(n_samples, n_samples, replace=True)
    return X[indices], y[indices]

def fit(self, X, y):
    self.n_classes_ = len(np.unique(y))
    self.n_features_ = X.shape[1]
    self.feature_importances_ = np.zeros(self.n_features_)

    for _ in range(self.n_estimators):
        # Generate a bootstrap sample
        X_sample, y_sample = self._bootstrap_sample(X, y)

        # Initialize a Decision Tree
        tree = DecisionTree(
            max_depth=self.max_depth,
            min_samples_split=self.min_samples_split,
            min_samples_leaf=self.min_samples_leaf
        )

        # Fit the tree on the bootstrap sample
        tree.fit(X_sample, y_sample)
        self.trees.append(tree)

    return self

def predict(self, X):
    # Get predictions from all trees
    predictions = np.array([tree.predict(X) for tree in self.trees])

    # Use majority voting to get final predictions
    return majority_vote(predictions)

```

Support Vector Machine Class

```

In [ ]: class SVM:
    def __init__(self, C=1.0, kernel='linear', max_iter=500, tol=1e-2, gamma='scale'):
        self.C = C
        self.kernel = kernel
        self.max_iter = max_iter
        self.tol = tol
        self.gamma = gamma

    def fit(self, X, y):
        n_samples = X.shape[0]
        self.X = X
        self.y = y

```

Support Vector Machine Class

```

In [ ]: class SVM:
    def __init__(self, C=1.0, kernel='linear', max_iter=500, tol=1e-2, gamma='scale'):
        self.C = C
        self.kernel = kernel
        self.max_iter = max_iter
        self.tol = tol
        self.gamma = gamma

    def fit(self, X, y):
        n_samples = X.shape[0]
        self.X = X
        self.y = y

        # Initialize parameters
        self.alphas = np.zeros(n_samples)
        self.b = 0
        self.K = self._compute_kernel_matrix(X)

        # Optimization Loop
        examine_all = True
        while True:
            num_changed = 0
            if examine_all:
                for i in range(n_samples):
                    E_i = self._compute_error(i)
                    if ((self.y[i] * E_i < -self.tol and self.alphas[i] < self.C) or
                        (self.y[i] * E_i > self.tol and self.alphas[i] > 0)):
                        num_changed += self._optimize_pair(i, E_i)
            else:
                non_bound_indices = np.where((self.alphas > 0) & (self.alphas < self.C))[0]
                for i in non_bound_indices:
                    E_i = self._compute_error(i)
                    if ((self.y[i] * E_i < -self.tol and self.alphas[i] < self.C) or
                        (self.y[i] * E_i > self.tol and self.alphas[i] > 0)):
                        num_changed += self._optimize_pair(i, E_i)

            if examine_all:
                examine_all = False
            elif num_changed == 0:
                examine_all = True

            if not examine_all and num_changed == 0:
                break

        return self

    def _compute_kernel_matrix(self, X):

```

```

def _compute_kernel_matrix(self, X):
    n_samples = X.shape[0]
    K = np.zeros((n_samples, n_samples))
    for i in range(n_samples):
        for j in range(n_samples):
            K[i,j] = self._kernel_function(X[i], X[j])
    return K

def _kernel_function(self, x1, x2):
    if self.kernel == 'linear':
        return np.dot(x1, x2)
    elif self.kernel == 'rbf':
        gamma = 1.0 / (len(x1) * np.var(x1)) if self.gamma == 'scale' else 1.0 / len(x1)
        return np.exp(-gamma * np.sum((x1 - x2) ** 2))
    raise ValueError(f"Unsupported kernel: {self.kernel}")

def _compute_error(self, i):
    return self._predict_raw(self.X[i]) - self.y[i]

def _optimize_pair(self, i, E_i):
    # Select second example
    j, E_j = self._select_second(i, E_i)

    # Skip if no improvement possible
    if i == j:
        return 0

    alpha_i_old = self.alphas[i]
    alpha_j_old = self.alphas[j]
    y_i = self.y[i]
    y_j = self.y[j]

    # Compute L and H
    if y_i != y_j:
        L = max(0, alpha_j_old - alpha_i_old)
        H = min(self.C, self.C + alpha_j_old - alpha_i_old)
    else:
        L = max(0, alpha_i_old + alpha_j_old - self.C)
        H = min(self.C, alpha_i_old + alpha_j_old)

    if L == H:
        return 0

    eta = 2 * self.K[i,j] - self.K[i,i] - self.K[j,j]
    if eta >= 0:
        return 0

    # Update alpha_j

```

```

# Update alpha_j
self.alphas[j] = alpha_j_old - (y_j * (E_i - E_j)) / eta
self.alphas[j] = min(H, max(L, self.alphas[j]))

if abs(self.alphas[j] - alpha_j_old) < 1e-5:
    return 0

# Update alpha_i
self.alphas[i] = alpha_i_old + y_i * y_j * (alpha_j_old - self.alphas[j])

# Update threshold
b1 = self.b - E_i - y_i * (self.alphas[i] - alpha_i_old) * self.K[i,i] \
    - y_j * (self.alphas[j] - alpha_j_old) * self.K[i,j]
b2 = self.b - E_j - y_i * (self.alphas[i] - alpha_i_old) * self.K[i,j] \
    - y_j * (self.alphas[j] - alpha_j_old) * self.K[j,j]

if 0 < self.alphas[i] < self.C:
    self.b = b1
elif 0 < self.alphas[j] < self.C:
    self.b = b2
else:
    self.b = (b1 + b2) / 2

return 1

def _select_second(self, i, E_i):
    # Select alpha j that maximizes the step size
    max_step = 0
    j = i
    E_j = 0

    valid_indices = np.where((self.alphas > 0) & (self.alphas < self.C))[0]
    if len(valid_indices) > 0:
        for k in valid_indices:
            if k == i:
                continue
            E_k = self._compute_error(k)
            step = abs(E_k - E_i)
            if step > max_step:
                max_step = step
                j = k
                E_j = E_k

    # If no non-bound examples found, choose random index
    if j == i:
        while j == i:
            j = np.random.randint(len(self.alphas))
        E_j = self._compute_error(j)

```

```

        j = np.random.randint(len(self.alphas))
        E_j = self._compute_error(j)

        return j, E_j

    def _predict_raw(self, x):
        result = 0
        for i in range(len(self.alphas)):
            if self.alphas[i] > 0:
                result += self.alphas[i] * self.y[i] * self._kernel_function(self.X[i], x)
        return result + self.b

    def predict(self, X):
        return np.array([np.sign(self._predict_raw(x)) for x in X])

class OneVsRestSVM:
    def __init__(self, C=1.0, kernel='linear', max_iter=500, tol=1e-2, gamma='scale'):
        self.C = C
        self.kernel = kernel
        self.max_iter = max_iter
        self.tol = tol
        self.gamma = gamma
        self.classifiers = {}

    def fit(self, X, y):
        self.classes_ = np.unique(y)

        for class_label in self.classes_:
            binary_y = np.where(y == class_label, 1, -1)
            classifier = SVM(
                C=self.C,
                kernel=self.kernel,
                max_iter=self.max_iter,
                tol=self.tol,
                gamma=self.gamma
            )
            classifier.fit(X, binary_y)
            self.classifiers[class_label] = classifier

        return self

    def predict(self, X):
        scores = np.zeros((X.shape[0], len(self.classes_)))
        for i, class_label in enumerate(self.classes_):
            classifier = self.classifiers[class_label]
            scores[:, i] = np.array([classifier._predict_raw(x) for x in X])
        return self.classes_[np.argmax(scores, axis=1)]

```

Model Evaluator Class

```

class ModelEvaluator:
    def __init__(self):
        pass

    def confusion_matrix(self, y_true: np.ndarray, y_pred: np.ndarray) -> np.ndarray:
        """
        Compute confusion matrix for multi-class classification.

        Args:
            y_true: Array of true labels
            y_pred: Array of predicted labels

        Returns:
            2D numpy array where element [i,j] is the number of observations with true label i that were predicted as label j
        """
        labels = sorted(list(set(np.concatenate([y_true, y_pred]))))
        n_labels = len(labels)
        label_to_index = {label: i for i, label in enumerate(labels)}

        # Initialize confusion matrix
        cm = np.zeros((n_labels, n_labels), dtype=int)

        # Fill confusion matrix
        for t, p in zip(y_true, y_pred):
            cm[label_to_index[t]][label_to_index[p]] += 1

        return cm, labels

    def accuracy_score(self, y_true: np.ndarray, y_pred: np.ndarray) -> float:
        """
        Compute classification accuracy.

        Args:
            y_true: Array of true labels
            y_pred: Array of predicted labels

        Returns:
            float: Accuracy score between 0.0 and 1.0
        """
        if len(y_true) != len(y_pred):
            raise ValueError("Length of true and predicted labels must match")

        return np.sum(y_true == y_pred) / len(y_true)

```

```

return np.sum(y_true == y_pred) / len(y_true)

def precision_recall_f1(self, y_true: np.ndarray, y_pred: np.ndarray) -> Dict[str, Dict[Union[str, float], float]]:
    """
    Compute precision, recall, and F1 score for each class.

    Args:
        y_true: Array of true labels
        y_pred: Array of predicted labels

    Returns:
        Dictionary containing precision, recall, and F1 scores for each class
    """
    cm, labels = self.confusion_matrix(y_true, y_pred)
    metrics = {}

    for idx, label in enumerate(labels):
        # True positives are diagonal elements
        tp = cm[idx, idx]
        # False positives are sum of column minus true positives
        fp = np.sum(cm[:, idx]) - tp
        # False negatives are sum of row minus true positives
        fn = np.sum(cm[idx, :]) - tp

        # Calculate metrics
        precision = tp / (tp + fp) if (tp + fp) > 0 else 0
        recall = tp / (tp + fn) if (tp + fn) > 0 else 0
        f1 = 2 * (precision * recall) / (precision + recall) if (precision + recall) > 0 else 0

        metrics[label] = {
            'precision': precision,
            'recall': recall,
            'f1': f1
        }

    # Calculate macro averaged metrics
    macro_precision = np.mean([m['precision'] for m in metrics.values()])
    macro_recall = np.mean([m['recall'] for m in metrics.values()])
    macro_f1 = np.mean([m['f1'] for m in metrics.values()])

    metrics['macro_avg'] = {
        'precision': macro_precision,
        'recall': macro_recall,
        'f1': macro_f1
    }

    return metrics

def classification_report(self, y_true: np.ndarray, y_pred: np.ndarray) -> str:
    """

```

```

def classification_report(self, y_true: np.ndarray, y_pred: np.ndarray) -> str:
    """
    Generate a formatted classification report including precision, recall, F1 score, and accuracy.

    Args:
        y_true: Array of true labels
        y_pred: Array of predicted labels

    Returns:
        str: Formatted classification report
    """
    metrics = self.precision_recall_f1(y_true, y_pred)
    accuracy = self.accuracy_score(y_true, y_pred)

    # Create formatted report
    report = ["Classification Report:", ""]
    report.append(f"{'Class':<15} {'Precision':>10} {'Recall':>10} {'F1-score':>10}")
    report.append("-" * 45)

    # Add metrics for each class
    for label in metrics:
        if label != 'macro_avg':
            m = metrics[label]
            report.append(f"{str(label):<15} {m['precision']:>10.4f} {m['recall']:>10.4f} {m['f1']:>10.4f}")

    # Add macro average
    report.append("-" * 45)
    macro = metrics['macro_avg']
    report.append(f"{'Macro avg':<15} {macro['precision']:>10.4f} {macro['recall']:>10.4f} {macro['f1']:>10.4f}")
    report.append(f"{'Accuracy':<15} {accuracy:>10.4f}")

    return "\n".join(report)

def ten_fold_cross_validation(self, model, X, y):
    """
    Perform 10-fold cross validation and return detailed metrics.

    Args:
        model: Model object with fit and predict methods
        X: Feature matrix
        y: Target labels

    Returns:
        Dictionary containing metrics for each fold
    """
    # Create 10 roughly equal-sized folds
    n_samples = len(X)
    fold_size = n_samples // 10
    indices = np.random.permutation(n_samples)

```

```

n_samples = len(x)
fold_size = n_samples // 10
indices = np.random.permutation(n_samples)

# Initialize lists to store metrics for each fold
fold_accuracies = []
fold_precisions = []
fold_recalls = []
fold_f1s = []

print("Performing 10-fold cross-validation...")

# Perform 10-fold cross validation
for fold in range(10):
    print(f"Processing fold {fold + 1}/10")

    # Create train/test split for this fold
    start_idx = fold * fold_size
    end_idx = start_idx + fold_size if fold < 9 else n_samples
    test_indices = indices[start_idx:end_idx]
    train_indices = np.concatenate([indices[:start_idx], indices[end_idx:]])

    X_train, X_test = X[train_indices], X[test_indices]
    y_train, y_test = y[train_indices], y[test_indices]

    # Train model
    model.fit(X_train, y_train)

    # Make predictions
    predictions = model.predict(X_test)

    # Calculate metrics
    metrics = self.precision_recall_f1(y_test, predictions)
    accuracy = self.accuracy_score(y_test, predictions)

    # Store metrics
    fold_accuracies.append(accuracy)
    fold_precisions.append(metrics['macro_avg']['precision'])
    fold_recalls.append(metrics['macro_avg']['recall'])
    fold_f1s.append(metrics['macro_avg']['f1'])

# Calculate average metrics across folds
avg_accuracy = np.mean(fold_accuracies)
avg_precision = np.mean(fold_precisions)
avg_recall = np.mean(fold_recalls)
avg_f1 = np.mean(fold_f1s)

# Calculate standard deviations
std_accuracy = np.std(fold_accuracies)
std_precision = np.std(fold_precisions)
std_recall = np.std(fold_recalls)
std_f1 = np.std(fold_f1s)

```

```

# Calculate standard deviations
std_accuracy = np.std(fold accuracies)
std_precision = np.std(fold precisions)
std_recall = np.std(fold recalls)
std_f1 = np.std(fold f1s)

print("\nCross-validation Results:")
print(f"Average Accuracy: {avg_accuracy:.4f} (+/- {std_accuracy:.4f})")
print(f"Average Precision: {avg_precision:.4f} (+/- {std_precision:.4f})")
print(f"Average Recall: {avg_recall:.4f} (+/- {std_recall:.4f})")
print(f"Average F1-score: {avg_f1:.4f} (+/- {std_f1:.4f})")

return {
    'accuracy': (avg_accuracy, std_accuracy),
    'precision': (avg_precision, std_precision),
    'recall': (avg_recall, std_recall),
    'f1': (avg_f1, std_f1)
}

```

Model Testing ¶

Dataset #1

```

# Load the dataset
file_path = r"C:\Users\aberg\OneDrive\Desktop\NJIT\Data Mining\Final Project\user_behavior_dataset.csv"
data = pd.read_csv(file_path)

# Display the first few rows of the dataset
print("Dataset Preview:")
display(data.head())

# Display basic information about the dataset
print("\nDataset Info:")
data.info()

# Check for missing values
print("\nMissing Values:")
print(data.isnull().sum())

# Display basic statistics to understand distributions
print("\nStatistical Summary:")
display(data.describe(include='all'))

```

```

# Data preprocessing
# Strip whitespace from column names
data.columns = data.columns.str.strip()

# Encode categorical features
data = one_hot_encode(data, 'Device Model')
data = one_hot_encode(data, 'Operating System')
data = one_hot_encode(data, 'Gender')

# Create target variable - keeping multiclass for this dataset
target_column = 'User Behavior Class'
y = data[target_column].values # Convert to numpy array
X = data.drop(columns=['User ID', target_column]).values # Convert to numpy array

# Standardize features
X = standardize(X)

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("\nTraining Random Forest...")
rf = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf.fit(X_train, y_train)
rf_predictions = rf.predict(X_test)

# SVM
print("\nTraining SVM...")
svm = OneVsRestSVM(C=1.0, kernel='linear', max_iter=200, tol=1e-3)
svm.fit(X_train, y_train)
svm_predictions = svm.predict(X_test)

# Evaluate both models
evaluator = ModelEvaluator()
print("\nRandom Forest Results:")
print(evaluator.classification_report(y_test, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test, svm_predictions))

print("\nEvaluating User Behavior Dataset:")
print("\nRandom Forest Results:")
rf_results = evaluator.ten_fold_cross_validation(rf, X, y)
print("\nSVM Results:")
svm_results = evaluator.ten_fold_cross_validation(svm, X, y)

```

Dataset #2

```

]: # Load the dataset
file_path = r"C:\Users\aberg\OneDrive\Desktop\NJIT\Data Mining\Final Project\loan_approval_dataset.csv"
data_loan = pd.read_csv(file_path)

# Display the first few rows of the dataset
print("Dataset Preview:")
display(data_loan.head())

# Display basic information about the dataset
print("\nDataset Info:")
data_loan.info()

# Check for missing values
print("\nMissing Values:")
print(data_loan.isnull().sum())

# Display basic statistics to understand distributions
print("\nStatistical Summary:")
display(data_loan.describe(include='all'))

]: # Use random 500 rows
data_loan = data_loan.sample(n=500, random_state=42)

# Strip whitespace from column names
data_loan.columns = data_loan.columns.str.strip()

# Create binary target (1 for Approved, 0 for Rejected)
data_loan['loan_status_binary'] = (data_loan['loan_status'] == 'Approved').astype(int)

# Encode categorical features
data_loan = one_hot_encode(data_loan, 'education')
data_loan = one_hot_encode(data_loan, 'self_employed')

# Separate features and target
X_loan = data_loan.drop(columns=['loan_id', 'loan_status', 'loan_status_binary']).values
y_loan = data_loan['loan_status_binary'].values

# Standardize features
X_loan = standardize(X_loan)

# Split data
X_train_loan, X_test_loan, y_train_loan, y_test_loan = train_test_split(X_loan, y_loan, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("Training Random Forest ...")

```

```

# Train and Evaluate both models
# Random Forest
print("Training Random Forest...")
rf_loan = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf_loan.fit(X_train_loan, y_train_loan)
rf_predictions = rf_loan.predict(X_test_loan)

# SVM
print("\nTraining SVM...")
svm_loan = OneVsRestSVM(C=1.0, kernel='linear', max_iter=200, tol=1e-3)
svm_loan.fit(X_train_loan, y_train_loan)
svm_predictions = svm_loan.predict(X_test_loan)

# Evaluate both models
evaluator = ModelEvaluator()
print("Random Forest Results:")
print(evaluator.classification_report(y_test_loan, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test_loan, svm_predictions))

```

```

print("\nEvaluating Loan Approval Dataset:")
print("\nRandom Forest Results:")
rf_results_loan = evaluator.ten_fold_cross_validation(rf_loan, X_loan, y_loan)
print("\nSVM Results:")
svm_results_loan = evaluator.ten_fold_cross_validation(svm_loan, X_loan, y_loan)

```

Dataset #3

```

# Load the dataset
file_path = r"C:\Users\aberg\OneDrive\Desktop\NJIT\Data Mining\Final Project\weather_classification_data.csv"
data_weather = pd.read_csv(file_path)

# Display the first few rows of the dataset
print("Dataset Preview:")
display(data_weather.head())

# Display basic information about the dataset
print("\nDataset Info:")
data_weather.info()

# Check for missing values
print("\nMissing Values:")
print(data_weather.isnull().sum())

# Display basic statistics to understand distributions
print("\nStatistical Summary:")
display(data_weather.describe(include='all'))

```

```

print("\nStatistical Summary:")
display(data_weather.describe(include='all'))

: # Use random 500 rows
data_weather = data_weather.sample(n=500, random_state=42)

# Fix the encoding and target creation
data_weather.columns = data_weather.columns.str.strip()
data_weather = one_hot_encode(data_weather, 'Cloud Cover')
data_weather = one_hot_encode(data_weather, 'Season')
data_weather = one_hot_encode(data_weather, 'Location')

# Create binary target for Sunny weather (multiclass classification)
target_column = 'Weather Type'
data_weather['target'] = (data_weather[target_column] == 'Sunny').astype(int)
data_weather = data_weather.drop(columns=[target_column])

# Separate features and target
X_weather = data_weather.drop(columns=['target']).values
y_weather = data_weather['target'].values

# Standardize features
X_weather = standardize(X_weather)

# Split data and continue with model training
X_train_weather, X_test_weather, y_train_weather, y_test_weather = train_test_split(X_weather, y_weather, test_size=0.3)

# Train and evaluate both models
# Random Forest
print("Training Random Forest...")
rf_weather = RandomForest(n_estimators=100, max_depth=10, min_samples_split=2, min_samples_leaf=1)
rf_weather.fit(X_train_weather, y_train_weather)
rf_predictions = rf_weather.predict(X_test_weather)

# SVM
print("\nTraining SVM...")
svm_weather = OneVsRestSVM(C=1.0, kernel='linear', max_iter=200, tol=1e-3)
svm_weather.fit(X_train_weather, y_train_weather)
svm_predictions = svm_weather.predict(X_test_weather)

# Evaluate
evaluator = ModelEvaluator()
print("\nRandom Forest Results:")
print(evaluator.classification_report(y_test_weather, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test_weather, svm_predictions))

svm_predictions = svm_weather.predict(X_test_weather)

# Evaluate
evaluator = ModelEvaluator()
print("\nRandom Forest Results:")
print(evaluator.classification_report(y_test_weather, rf_predictions))
print("\nSVM Results:")
print(evaluator.classification_report(y_test_weather, svm_predictions))

print("\nEvaluating Weather Dataset:")
print("\nRandom Forest Results:")
rf_results_weather = evaluator.ten_fold_cross_validation(rf_weather, X_weather, y_weather)
print("\nSVM Results:")
svm_results_weather = evaluator.ten_fold_cross_validation(svm_weather, X_weather, y_weather)

```

9.4 References

Khorasani, Val. *Mobile Device Usage and User Behavior Dataset*. 2023. Kaggle, <https://www.kaggle.com/datasets/valakhorasani/mobile-device-usage-and-user-behavior-dataset>.

Sharma, Archit. *Loan Approval Prediction Dataset*. 2023. Kaggle, <https://www.kaggle.com/datasets/architsharma01/loan-approval-prediction-dataset>.

Nikhil7280. *Weather Type Classification*. 2023. Kaggle, <https://www.kaggle.com/datasets/nikhil7280/weather-type-classification>.

Analytics Vidhya. *Building a Random Forest from Scratch: Understanding Real-World Data Products – ML for Programmers (Part 3)*. 2018. Analytics Vidhya, <https://www.analyticsvidhya.com/blog/2018/12/building-a-random-forest-from-scratch-understanding-real-world-data-products-ml-for-programmers-part-3/>.

Brownlee, Jason. *How to Implement Random Forest From Scratch in Python*. Machine Learning Mastery, 7 April 2021, <https://machinelearningmastery.com/implement-random-forest-scratch-python/>.

Ninjaac. *Random Forest From Scratch*. Kaggle, <https://www.kaggle.com/code/ninjaac/random-forest-from-scratch>.

PyCodeMates. *Implementing SVM From Scratch in Python*. 2022. PyCodeMates, <https://www.pycodemates.com/2022/10/implementing-SVM-from-scratch-in-python.html>.

Python Engineer. *SVM From Scratch*. Python Engineer, https://www.python-engineer.com/courses/mlfromscratch/07_svm/.

Jha, Himanshu. *Implement Multiclass SVM From Scratch in Python*. Towards Data Science, 6 May 2021, <https://towardsdatascience.com/implement-multiclass-svm-from-scratch-in-python-b141e43dc084>.

Plain English. *Support Vector Machines (SVM) From Scratch to Libraries*. Plain English, 13 January 2023, <https://python.plainenglish.io/support-vector-machines-svm-from-scratch-to-libraries-990086a666d4>.

The Pandas Development Team. *Pandas: Python Data Analysis Library*. <https://pandas.pydata.org/>.

The NumPy Development Team. *NumPy*. <https://numpy.org/>.

Python Software Foundation. *Typing Module*. Python 3 Documentation, <https://docs.python.org/3/library/typing.html>.

OpenAI. *ChatGPT: AI Language Model*. OpenAI, 2024, <https://openai.com/chatgpt>. (Generative AI was used solely for the purpose of debugging, explaining, and documenting code/algorithms used in the development of this project. Also used for the purpose of project management and assistance)